

Data Wrangling

August 26, 2023

```
# **
```

Data Wrangling**

Import Packages

```
[99]: import os
      from dotenv import load_dotenv

      import datetime as dt

      from tqdm import tqdm

      import math
      import numpy as np
      import pandas as pd

      import matplotlib.pyplot as plt
      import seaborn as sns

      from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import Lasso
      from sklearn.model_selection import GridSearchCV

      from statsmodels.stats.outliers_influence import variance_inflation_factor
```

Set Correct Options for the Notebook

```
[100]: # Load in the dotenv variables
      load_dotenv()
      # Show all the columns in the .head() method
      pd.set_option('display.max_columns', None)
      # Turn all plots dark to help my eyes
      plt.style.use('dark_background')
```

The data was taken from kaggle and contains [Loan Data](#) information on the borrowers. Below I used Kaggle's API to download the dataset. I had trouble getting it to work inline in the notebook so I used these commands in the terminal.

```
[101]: '''
# Download the data from the kaggle api
kaggle datasets download -d ranadeep/credit-risk-dataset
# unzip the data
unzip credit-risk-dataset.zip -d credit-risk-dataset
'''
```

```
[101]: '\n# Download the data from the kaggle api\nkaggle datasets download -d
ranadeep/credit-risk-dataset\n# unzip the data \nunzip credit-risk-dataset.zip
-d credit-risk-dataset \n'
```

Set Directory

```
[102]: # Get the path from the environment from the dotenv file without extra add-ons
project_path = os.getenv('Project_Path')[2:78]

# Change notebook directory back one so that it can access the data
os.chdir(project_path)
```

Load the Data

```
[103]: # Load in the data with low_memory as False so that it can understand dtype
loan = pd.read_csv('./data/raw/loan.csv', low_memory = False)
# Print out the shape of the dataset
print(loan.shape)
# Print out a sample of the dataset
loan.head()
```

(887379, 74)

```
[103]:
```

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	\
0	1077501	1296599	5000.0	5000.0	4975.0	36 months	
1	1077430	1314167	2500.0	2500.0	2500.0	60 months	
2	1077175	1313524	2400.0	2400.0	2400.0	36 months	
3	1076863	1277178	10000.0	10000.0	10000.0	36 months	
4	1075358	1311748	3000.0	3000.0	3000.0	60 months	

	int_rate	installment	grade	sub_grade	emp_title	emp_length	\
0	10.65	162.87	B	B2	NaN	10+ years	
1	15.27	59.83	C	C4	Ryder	< 1 year	
2	15.96	84.33	C	C5	NaN	10+ years	
3	13.49	339.31	C	C1	AIR RESOURCES BOARD	10+ years	
4	12.69	67.79	B	B5	University Medical Group	1 year	

	home_ownership	annual_inc	verification_status	issue_d	loan_status	\
0	RENT	24000.0	Verified	Dec-2011	Fully Paid	
1	RENT	30000.0	Source Verified	Dec-2011	Charged Off	
2	RENT	12252.0	Not Verified	Dec-2011	Fully Paid	
3	RENT	49200.0	Source Verified	Dec-2011	Fully Paid	

4	RENT	80000.0	Source Verified	Dec-2011	Current
---	------	---------	-----------------	----------	---------

	pymnt_plan	url \
0	n	https://www.lendingclub.com/browse/loanDetail...
1	n	https://www.lendingclub.com/browse/loanDetail...
2	n	https://www.lendingclub.com/browse/loanDetail...
3	n	https://www.lendingclub.com/browse/loanDetail...
4	n	https://www.lendingclub.com/browse/loanDetail...

	desc	purpose \
0	Borrower added on 12/22/11 > I need to upgra...	credit_card
1	Borrower added on 12/22/11 > I plan to use t...	car
2		NaN small_business
3	Borrower added on 12/21/11 > to pay for prop...	other
4	Borrower added on 12/21/11 > I plan on combi...	other

	title	zip_code	addr_state	dti	delinq_2yrs \
0	Computer	860xx	AZ	27.65	0.0
1	bike	309xx	GA	1.00	0.0
2	real estate business	606xx	IL	8.72	0.0
3	personel	917xx	CA	20.00	0.0
4	Personal	972xx	OR	17.94	0.0

	earliest_cr_line	inq_last_6mths	mths_since_last_delinq \
0	Jan-1985	1.0	NaN
1	Apr-1999	5.0	NaN
2	Nov-2001	2.0	NaN
3	Feb-1996	1.0	35.0
4	Jan-1996	0.0	38.0

	mths_since_last_record	open_acc	pub_rec	revol_bal	revol_util \
0	NaN	3.0	0.0	13648.0	83.7
1	NaN	3.0	0.0	1687.0	9.4
2	NaN	2.0	0.0	2956.0	98.5
3	NaN	10.0	0.0	5598.0	21.0
4	NaN	15.0	0.0	27783.0	53.9

	total_acc	initial_list_status	out_prncp	out_prncp_inv	total_pymnt \
0	9.0	f	0.0	0.0	5861.071414
1	4.0	f	0.0	0.0	1008.710000
2	10.0	f	0.0	0.0	3003.653644
3	37.0	f	0.0	0.0	12226.302212
4	38.0	f	766.9	766.9	3242.170000

	total_pymnt_inv	total_rec_prncp	total_rec_int	total_rec_late_fee \
0	5831.78	5000.00	861.07	0.00
1	1008.71	456.46	435.17	0.00

2	3003.65	2400.00	603.65	0.00
3	12226.30	10000.00	2209.33	16.97
4	3242.17	2233.10	1009.07	0.00

	recoveries	collection_recovery_fee	last_pymnt_d	last_pymnt_amnt	\
0	0.00	0.00	Jan-2015	171.62	
1	117.08	1.11	Apr-2013	119.66	
2	0.00	0.00	Jun-2014	649.91	
3	0.00	0.00	Jan-2015	357.48	
4	0.00	0.00	Jan-2016	67.79	

	next_pymnt_d	last_credit_pull_d	collections_12_mths_ex_med	\
0	NaN	Jan-2016	0.0	
1	NaN	Sep-2013	0.0	
2	NaN	Jan-2016	0.0	
3	NaN	Jan-2015	0.0	
4	Feb-2016	Jan-2016	0.0	

	mths_since_last_major_derog	policy_code	application_type	\
0	NaN	1.0	INDIVIDUAL	
1	NaN	1.0	INDIVIDUAL	
2	NaN	1.0	INDIVIDUAL	
3	NaN	1.0	INDIVIDUAL	
4	NaN	1.0	INDIVIDUAL	

	annual_inc_joint	dti_joint	verification_status_joint	acc_now_delinq	\
0	NaN	NaN	NaN	0.0	
1	NaN	NaN	NaN	0.0	
2	NaN	NaN	NaN	0.0	
3	NaN	NaN	NaN	0.0	
4	NaN	NaN	NaN	0.0	

	tot_coll_amt	tot_cur_bal	open_acc_6m	open_il_6m	open_il_12m	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	

	open_il_24m	mths_since_rcnt_il	total_bal_il	il_util	open_rv_12m	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	

	open_rv_24m	max_bal_bc	all_util	total_rev_hi_lim	inq_fi	total_cu_tl	\
--	-------------	------------	----------	------------------	--------	-------------	---

0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN

	inq_last_12m
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

The data is large and complex. In the `read_csv()` function I had to specify `low_memory = False` because some columns had too many different data types for the `read_csv()` function to be able to quickly determine the data type of each column.

The first step in data wrangling is to separate out the features (X) from the target variable (y). When I initially started this project I had believed that y was `loan_status` because it was listed on the kaggle website, and was a clear representation of how the loan had fared. After much data wrangling and looking closely at the dataset I realized that it was hard to use many of the features while having `loan_status` be the target variable. Too many of the features have information on the loan after origination. If the purpose of the project was to try to find the best loans to invest in, then I cannot use features that are collected during the term of the loan.

In order to use many of the features of the dataset, I need a question that has a later time frame. If I try to find how much each borrower will still pay back on their loans, I can use the features that are during the term of the loan. This can create a market or evaluation for the value of the existing loans that would be valuable for loan investors or institutions who are thinking of selling off their loans for immediate liquidity.

Below I separated out every column of the dataset that had information on the amount that the borrower has already paid and the amount that they owe. With the remaining columns I'll be able to create a multiplier that will provide an estimate on how much the borrower will pay.

```
[104]: # Create the features dataframe
X = loan.drop(['out_prncp', 'out_prncp_inv', 'total_pymnt', 'total_pymnt_inv',
              'total_rec_prncp', 'total_rec_int',
              'total_rec_late_fee', 'recoveries',
              'collection_recovery_fee', 'loan_status'], axis = 1)
# Separate out the target variable
y = loan[['out_prncp', 'out_prncp_inv', 'total_pymnt', 'total_pymnt_inv',
          'total_rec_prncp',
          'total_rec_int', 'total_rec_late_fee', 'recoveries',
          'collection_recovery_fee']]
```

0.1 Cleaning up the Dataset

Unique Values The easiest features to clean are ones that have constant values since they provide no information. Below I looked at how many unique values each feature has. I will drop every feature that has only 1 unique value.

```
[105]: X.nunique().sort_values(ascending = True)
```

```
[105]: policy_code          1
      pymnt_plan          2
      application_type    2
      initial_list_status  2
      term                2
      ...
      emp_title          299271
      tot_cur_bal        327342
      url                887379
      member_id          887379
      id                 887379
      Length: 64, dtype: int64
```

```
[106]: # Drop the 'policy_code' feature since there is only one policy code
      X.drop('policy_code', axis = 1, inplace = True)
```

0.1.1 Categorical Columns

In order to analyze the data, all of the features have to have only numerical values. Some columns contain numbers but need to be cleaned in order to be fully numeric where as others are purely categorical and either need to be one-hot encoded or dropped if they have too many unique values.

Below I printed out a list of all of the categorical columns.

```
[107]: # Use dtypes to find all of the columns that are not int or float in X
      [X.dtypes.index[i] for i,type in enumerate(X.dtypes) if type not in_
      ↪('int64','float64','uint8')]
```

```
[107]: ['term',
      'grade',
      'sub_grade',
      'emp_title',
      'emp_length',
      'home_ownership',
      'verification_status',
      'issue_d',
      'pymnt_plan',
      'url',
      'desc',
      'purpose',
      'title',
```

```
'zip_code',
'addr_state',
'earliest_cr_line',
'initial_list_status',
'last_pymnt_d',
'next_pymnt_d',
'last_credit_pull_d',
'application_type',
'verification_status_joint']
```

term of the loan

```
[108]: X['term'].unique()
```

```
[108]: array([' 36 months', ' 60 months'], dtype=object)
```

The **term** variable represents how long loan is for and has two values: 36 months and 60 months. I can convert those strings to 3, and 5 respectively to represent the years.

```
[109]: # replace the string values with numerical values, use a dictionary to be
↳concise
X['term'].replace({' 36 months':3, ' 60 months':5}, inplace = True)
# Show the column
X['term']
```

```
[109]: 0      3
      1      5
      2      3
      3      3
      4      5
      ..
      887374  3
      887375  3
      887376  5
      887377  5
      887378  3
      Name: term, Length: 887379, dtype: int64
```

grade and sub_grade for the loans

```
[110]: # Show all the different possible grades
print('Grade:', X['grade'].unique())

# Show all the different possible sub grades
print('Grade:', X['sub_grade'].unique())
```

```
Grade: ['B' 'C' 'A' 'E' 'F' 'D' 'G']
Grade: ['B2' 'C4' 'C5' 'C1' 'B5' 'A4' 'E1' 'F2' 'C3' 'B1' 'D1' 'A1' 'B3' 'B4'
       'C2' 'D2' 'A3' 'A5' 'D5' 'A2' 'E4' 'D3' 'D4' 'F3' 'E3' 'F4' 'F1' 'E5'
       'G4' 'E2' 'G3' 'G2' 'G1' 'F5' 'G5']
```

There is also the subgrade column which differentiates the quality of loans within the different grades. Since there are 5 subgrades for each grade, I can convert the subgrade to .2 downgrade of the grade

```
[111]: # Convert the grades into numbers that can be analyzed
X['grade'] = X['grade'].replace({'A':7, 'B':6, 'C':5, 'D':4, 'E':3, 'F':2, 'G':
    ↪1})
# Add in the subgrade as a .2 change in the grade
X['grade'] = [X['grade'][i] - ((int(sub[-1]) - 1)/5) for i, sub in
    ↪enumerate(X['sub_grade'])]
# Show the grade
X['grade']
```

```
[111]: 0          5.8
      1          4.4
      2          4.2
      3          5.0
      4          5.2
      ...
      887374      5.2
      887375      5.2
      887376      3.8
      887377      2.6
      887378      5.2
      Name: grade, Length: 887379, dtype: float64
```

With grade incorporating both grade and sub_grade, I can drop sub_grade.

```
[112]: # Drop sub_grade from X
X.drop('sub_grade', axis = 1, inplace = True)
```

emp_title: The type of employment for the borrower.

```
[113]: X['emp_title'].unique()
```

```
[113]: array([nan, 'Ryder', 'AIR RESOURCES BOARD', ..., 'machining Cell Lead',
      'KYC Business Analyst', 'Manager Hotel Operations Oasis '],
      dtype=object)
```

Unfortunately it appears that many of the borrowers put the name of their employer instead of their profession. I can confirm this better by using the `value_counts()` method from pandas.

```
[114]: X['emp_title'].value_counts()
```

```
[114]: Teacher          13469
      Manager          11240
      Registered Nurse    5525
      Owner            5376
      RN              5355
```



```

Thomas J. Paul, Inc.      ...      1
Piggie Toes Preschool    ...      1
greystone park psychiatric hospital    1
Las Vegas Motropolitan Police Department    1
Manager Hotel Operations Oasis    1
Name: emp_title, Length: 299271, dtype: int64

```

While the most common values are types of professions, it appears that there are many values that have just the name of the employers. If I take a look at all of the values in `emp_title` that appear over 1000 times I'll get a better sense of the common professions I should try to simplify the column into.

```

[115]: # find all the different types of jobs that the borrower could have
print(len(X['emp_title'].value_counts()[X['emp_title'].value_counts() > 1000]))
# See how common each of the different employment types
X['emp_title'].value_counts()[X['emp_title'].value_counts() > 1000]

```

48

```

[115]: Teacher      13469
Manager      11240
Registered Nurse      5525
Owner      5376
RN      5355
Supervisor      4983
Sales      4212
Project Manager      3988
Driver      3569
Office Manager      3510
General Manager      3178
Director      3156
manager      3138
teacher      2925
owner      2849
Engineer      2671
President      2598
driver      2429
Vice President      2351
Attorney      2136
Operations Manager      2071
Accountant      2035
Administrative Assistant      2019
Sales Manager      1846
Account Manager      1725
sales      1724
Police Officer      1720
supervisor      1675

```

Executive Assistant	1603
Analyst	1538
Store Manager	1515
Technician	1462
Nurse	1426
truck driver	1387
Truck Driver	1387
Software Engineer	1344
Paralegal	1299
Controller	1246
Consultant	1242
Assistant Manager	1199
Program Manager	1188
Branch Manager	1177
Server	1120
Administrator	1103
Principal	1073
Account Executive	1072
Mechanic	1044
Business Analyst	1013

Name: emp_title, dtype: int64

What stands out in this column is that many of the employment types seem to be variations on the same profession. For instance there are a lot of different kinds of managers even though they are all in the managerial profession. I can combine all of these to create a much more consolidated column.

Below I created a dictionary where the keys are the profession and the values are lists of all phrases that would be associated with that profession.

```
[116]: Professions = {'Manager':['manag','supervisor','superintendent', ' lead ',␣
    ↪'foreman'],
    'Healer':['nurse', 'nursing', 'cna', 'lpn', 'physician',␣
    ↪'doctor', 'pharmacist', 'counselor', 'therapist', 'rn', 'paramedic'],
    'Technical': ['engineer','mechanic', 'electrician', 'machinist',␣
    ↪'machining', 'technician', 'it', 'software', 'tech','welder'],
    'Executive':['president', 'owner', 'ceo', 'partner', 'vp'],␣
    ↪'Designer': ['designer'],
    'Vol': ['firefighter', 'social worker', 'army', 'officer',␣
    ↪'sheriff', 'deputy', 'sergeant', 'agent', 'colonel','USAF'],
    'Director':['director', 'coordinator'], 'Accountant':
    ↪['accountant', 'bookkeeper', 'controller','accounting'], 'Sales':['sales',␣
    ↪'realt'],
    'Finance':['financial','underwriter', 'broker', 'cfo','bank'],␣
    ↪'Analyst':['analyst','consult', 'analysis','estimator'],
    'Clergy':['pastor','rabbi','priest', 'imam', 'minister'],␣
    ↪'Clerk':['clerk'],
```

```

        'Service':['bartender', 'server', 'service', 'diner', 'grill',
↪ 'attendant', 'cashier', 'aesthetician', 'dealer', 'cook', 'chef'],
        'Manlab':['maintenance', 'maintenence', 'laborer', 'lumber',
↪ 'custodian', 'carpenter'],
        'Operator':['operator', 'pilot', 'driver'], 'Assistant':
↪ ['assistant', 'secretary', 'receptionist'], 'Law':
↪ ['attorney', 'esq', 'paralegal'],
        'Education':
↪ ['teacher', 'principal', 'professor', 'school', 'educator'], 'Admin':['admin']}

```

Below I created a new column that assumed that all missing values were unemployed. Then I iterated through the Professions dictionary and for each value I searched through the entire column to see if any cells contained that value. If they did, then I changed the entire cell to be the corresponding key from the dictionary.

```

[117]: # Fill in the missing values with unemployed for analysis
X['emp_title_cons'] = X['emp_title'].fillna('Unemployed')
for key, value in tqdm(Professions.items()):
    for v in value:
        # Create a boolean mask for all employments that are managerial
        prof = X['emp_title_cons'].str.contains(v, case=False)
        # Change the value of those to 'Manager'
        X.loc[prof, 'emp_title_cons'] = key

```

```

0%|          | 0/20 [00:00<?, ?it/s]100%|          | 20/20 [00:32<00:00,
1.65s/it]

```

With the employment titles consolidated, I can look at the employment titles with over 1000 instances again to see how effective the grouping was.

```

[118]: # find all the different types of jobs that the borrower could have
print(len(X['emp_title_cons'].value_counts()[X['emp_title_cons'].value_counts()
↪ > 1000]))

# See how common each of the different employment types
X['emp_title_cons'].value_counts()[X['emp_title_cons'].value_counts() > 1000]

```

22

```

[118]: Manager          141239
      Technical        119770
      Unemployed       51462
      Healer           50322
      Director         32973
      Education        29821
      Executive         28111
      Vol              25647
      Operator         24758
      Service          22252

```

Analyst	21481
Assistant	18877
Sales	17653
Finance	12007
Accountant	10458
Admin	9648
Manlab	5969
Clerk	5333
Designer	2466
Law	1942
Clergy	1392
Account Executive	1072

Name: emp_title_cons, dtype: int64

There is a marked improvement here with “Manager” accounting for over 150,000 of the observations. It is possible that some of the profession types are mistakes but I intentionally put spaces before and after ‘lead’ so that the code would only capture instances of the word lead itself.

Below I changed the `emp_type` column to have only professions with more than 1000 instances. The rest I classified as other.

```
[119]: common_proffs = X['emp_title_cons'].value_counts()[X['emp_title_cons'].
        ↪value_counts() > 1000]
        # Create a feature that has the most common profession types, store the rest as
        ↪other.
X['emp_type'] = [x if x in common_proffs.index else 'other' for x in
        ↪X['emp_title_cons']]
```

With a workable list of professions that the borrowers have, I can one-hot encode them for later analysis and drop the columns that created `emp_type`.

```
[120]: # Convert all of the columns into dummy variables
X = pd.get_dummies(X, columns = ['emp_type'], drop_first = True)

        # Drop emp_title and emp_title_cons now that there is no need for them.
X.drop(['emp_title', 'emp_title_cons'], axis = 1, inplace = True)
```

emp_length: employment length of the borrower.

```
[121]: # Show all of the different employment lengths
X['emp_length'].unique()
```

```
[121]: array(['10+ years', '< 1 year', '1 year', '3 years', '8 years', '9 years',
        '4 years', '5 years', '6 years', '2 years', '7 years', nan],
        dtype=object)
```

This column appears to have numerical values that are encased in strings. If I can separate the number from the surrounding string and change its type to a number then this column can become numerical.

```
[122]: # Replace the non-existent values with 0
X['emp_length'] = X['emp_length'].fillna(0)
# Replace the less than one year with the average of .5
X['emp_length'] = X['emp_length'].replace('< 1 year', 0.5)
# Get rid of all non-numerical characters to convert the values into numerical
X['emp_length'] = X['emp_length'].replace(r'[a-zA-Z+]', '', regex = True)
# Change the type from string to numerical
X['emp_length'] = X['emp_length'].astype('float')
```

home_ownership: The housing situation of the borrower

```
[123]: # Look at all the types of home ownership
X['home_ownership'].unique()
```

```
[123]: array(['RENT', 'OWN', 'MORTGAGE', 'OTHER', 'NONE', 'ANY'], dtype=object)
```

Since OTHER, NONE, and ANY can all mean the same thing, I think it best to combine the three of them into OTHER

```
[124]: # Switch home_ownership rare values to other.
X['home_ownership'].replace({'NONE': 'OTHER', 'ANY': 'OTHER'}, inplace = True)
```

With a consolidated home_ownership column, I can now one-hot encode it for later analysis

```
[125]: # Create dummy columns for the different types of home ownership
X = pd.get_dummies(X, columns = ['home_ownership'], drop_first = True)
```

verification_status: Whether or not Lending club was able to verify the information the borrower submitted.

```
[126]: X['verification_status'].unique()
```

```
[126]: array(['Verified', 'Source Verified', 'Not Verified'], dtype=object)
```

verification_status has 3 unique values that determine if the information has been verified, the source of the information has been verified or if neither the source nor the information has been verified. I will make this into a dummy variable later on.

```
[127]: # Create dummy variables for the different verification statuses
X = pd.get_dummies(X, columns = ['verification_status'], drop_first = True)
```

issue_d: the month and year in which the loan was issued.

```
[128]: X['issue_d'].head()
```

```
[128]: 0    Dec-2011
1    Dec-2011
2    Dec-2011
3    Dec-2011
4    Dec-2011
Name: issue_d, dtype: object
```

`issue_d` is a very important feature because I'll be able to use it to determine how much time is left on the loan but it isn't a feature that will influence a borrower's ability to pay back the loan. For that reason I'll drop it from this dataset.

```
[129]: # Drop issue_d from the X dataset
X.drop('issue_d', axis = 1, inplace = True)
```

`pymnt_plan`: Indicates whether there has been a payment plan put in place for the borrower to pay off the loan that wasn't the original schedule.

```
[130]: X['pymnt_plan'].unique()
```

```
[130]: array(['n', 'y'], dtype=object)
```

There are only two values in this column either a 'n' for no or a 'y' for yes. Since it is a binary choice I can represent it numerically as 0 for n and 1 for y.

```
[131]: # Convert pymnt_plan to a binary variable
X['pymnt_plan'] = np.where(X['pymnt_plan'] == 'y', 1, 0)
```

`url` of the page that contains the loan. Needs a login to access it.

```
[132]: X['url'].unique()
```

```
[132]: array(['https://www.lendingclub.com/browse/loanDetail.action?loan_id=1077501',
          'https://www.lendingclub.com/browse/loanDetail.action?loan_id=1077430',
          'https://www.lendingclub.com/browse/loanDetail.action?loan_id=1077175',
          ...,
          'https://www.lendingclub.com/browse/loanDetail.action?loan_id=36271333',
          'https://www.lendingclub.com/browse/loanDetail.action?loan_id=36490806',
          'https://www.lendingclub.com/browse/loanDetail.action?loan_id=36271262'],
          dtype=object)
```

The URL column appears to be the same website with the id column as the `loan_id` at the end of it. I can test this by trying to see if the characters up until the `id=` are always the same.

```
[133]: # Grab the first part of the url and see if it always the same.
X['url'].str[:61].unique()
```

```
[133]: array(['https://www.lendingclub.com/browse/loanDetail.action?loan_id='],
          dtype=object)
```

Now that I know that the start of the `url` is always the same I can check to see if the ID part of the `url` is the exact same as the `id` column itself.

```
[134]: # subtract the set of the id part of the url column from the id column.
set(X['id']) - set(X['url'].str[61:].astype('int'))
```

```
[134]: set()
```

The code above returned an empty set which means that there is no difference between the id column and the number at the end of the url. This means that the url is redundant and can be dropped.

```
[135]: # Drop the url column from X
X.drop('url', axis = 1, inplace = True)
```

desc: A summary for the reason the borrower is taking the loan

```
[136]: # Find how many people didn't put a description
print(X['desc'].isna().sum())

# Print a sample of some of the descriptions
X['desc'].unique()
```

761351

```
[136]: array([' Borrower added on 12/22/11 > I need to upgrade my business
technologies.<br>',
      ' Borrower added on 12/22/11 > I plan to use this money to finance the
motorcycle i am looking at. I plan to have it paid off as soon as possible/when
i sell my old bike. I only need this money because the deal im looking at is to
good to pass up.<br><br> Borrower added on 12/22/11 > I plan to use this money
to finance the motorcycle i am looking at. I plan to have it paid off as soon as
possible/when i sell my old bike.I only need this money because the deal im
looking at is to good to pass up. I have finished college with an associates
degree in business and its takingmeplaces<br>',
      nan, ...,
      'I need a lower interest loan to pay off my citifinancial loan. ',
      'I am looking for a loan to pay my credit cards off as well as making
some very much needed auto repairs',
      'I am in my senior year of college in obtaining a bachelors degree in
criminal justice. I do not qualify for financial aid and have used all stafford
loans available. My tuition is approx. $1200 a month and I have 10 courses left
which adds up to over $10,000. I need some assistance to cover my tuition until
my graduation date which is August of 2009.'],
      dtype=object)
```

The desc variable is significantly more in depth than the title or purpose variable but unfortunately most of the borrowers never submitted one. Since most of the important information that could be gleaned from the desc column is most likely contained within title and purpose which have little to no missing values, I will drop this column.

```
[137]: X.drop('desc', axis = 1, inplace = True)
```

purpose: The reason the borrower applied and took the loan

```
[138]: # View all the unique reasons for why people are getting a loan.
X['purpose'].unique()
```

```
[138]: array(['credit_card', 'car', 'small_business', 'other', 'wedding',
        'debt_consolidation', 'home_improvement', 'major_purchase',
        'medical', 'moving', 'vacation', 'house', 'renewable_energy',
        'educational'], dtype=object)
```

This column needs to be one-hot encoded since there are no clear numerical map for the categorical values of this column. .

```
[139]: X = pd.get_dummies(X, columns = ['purpose'], drop_first = True)
```

title: The title that the borrower put for the loan

```
[140]: print(len(X['title'].unique()))
        X['title'].unique()
```

63145

```
[140]: array(['Computer', 'bike', 'real estate business', ...,
        'new kitchen for momma!',
        'New Baby and New House (CC Consolidate)',
        'Credit Card/Auto Repair'], dtype=object)
```

This appears to be a worse version of purpose as there are over 63,000 different titles and the content from them appears to be very similar to the content in the `purpose` variable. For that reason I'll drop title

```
[141]: X.drop('title', axis = 1, inplace = True)
```

zip_code: The borrower's zipcode

```
[142]: print(len(X['zip_code'].unique()))
        X['zip_code'].head()
```

935

```
[142]: 0    860xx
        1    309xx
        2    606xx
        3    917xx
        4    972xx
        Name: zip_code, dtype: object
```

While it looks like there are numbers to be used in this column, these are zipcodes and not orders of magnitude. Therefore if I removed the xs I would still have a categorical column. I could one-hot encode the zip codes but that would produce over 800 million values and I'm not sure that my laptop would be able to easily run calculations on that amount of data so the simplest thing to do is to drop them.

```
[143]: # Drop zip_code from the X dataset
        X.drop('zip_code', axis = 1, inplace = True)
```

addr_state: Initials for the state the borrower is from


```
[144]: # Print out the number of states that are used
print(len(X['addr_state'].unique()))
# Print out the list of states used.
(X['addr_state'].unique())
```

51

```
[144]: array(['AZ', 'GA', 'IL', 'CA', 'OR', 'NC', 'TX', 'VA', 'MO', 'CT', 'UT',
        'FL', 'NY', 'PA', 'MN', 'NJ', 'KY', 'OH', 'SC', 'RI', 'LA', 'MA',
        'WA', 'WI', 'AL', 'CO', 'KS', 'NV', 'AK', 'MD', 'WV', 'VT', 'MI',
        'DC', 'SD', 'NH', 'AR', 'NM', 'MT', 'HI', 'WY', 'OK', 'DE', 'MS',
        'TN', 'IA', 'NE', 'ID', 'IN', 'ME', 'ND'], dtype=object)
```

While there are 50 states, there are borrowers from Washington, D.C. brings the total to 51. Unlike zipcodes, 51 is a much more reasonable number to one-hot encode.

```
[145]: # Create a dummy variable for each state that a borrower is from
X = pd.get_dummies(X, columns = ['addr_state'], drop_first = True)
```

earliest_cr_line: The date of the earliest line of credit that the borrower had

```
[146]: print(X['earliest_cr_line'].isna().sum())
print(X['earliest_cr_line'].dtype)
X['earliest_cr_line'].head()
```

29

object

```
[146]: 0    Jan-1985
1    Apr-1999
2    Nov-2001
3    Feb-1996
4    Jan-1996
Name: earliest_cr_line, dtype: object
```

While it does appear that this column is in a datetime format, it is actually an object that needs to be converted. There are still 29 missing values but I will fill them in with the `issue_d` as the loan itself is a form of credit.

```
[147]: # Fill in missing dates for the earliest credit line column
X['earliest_cr_line'] = X['earliest_cr_line'].fillna(loan['issue_d'])

# Convert the earliest credit line of the borrower
X['earliest_cr_line'] = pd.to_datetime(X['earliest_cr_line'])
```

In order to include this variable in the multicollinearity and lasso calculations I need to convert it from datetime to numeric.

```
[148]: # Convert the variable from datetime to numeric.
X['earliest_cr_line'] = X['earliest_cr_line'].map(dt.datetime.toordinal)
```

initial_list_status: Listing status for if the loan was offered as whole or fractional.

```
[149]: X['initial_list_status'].unique()
```

```
[149]: array(['f', 'w'], dtype=object)
```

‘w’ represents loans that were offered to institutions that would be bought as a full loan. ‘f’ represents loans that went straight to investors who would buy a portion of the loan instead of all of it. Since there are only two values, I can binary encode this variable to have ‘f’ → 0 and ‘w’ → 1.

```
[150]: # Convert pymnt_plan to a binary variable
X['initial_list_status'] = np.where(X['initial_list_status'] == 'w', 1, 0)
```

last_pymnt_d: The last time that the borrower paid lending club in mm-yyyy

```
[151]: print(X['last_pymnt_d'].isna().sum())
X['last_pymnt_d'].head(10)
```

17659

```
[151]: 0    Jan-2015
1    Apr-2013
2    Jun-2014
3    Jan-2015
4    Jan-2016
5    Jan-2015
6    Jan-2016
7    Jan-2015
8    Apr-2012
9    Nov-2012
Name: last_pymnt_d, dtype: object
```

This variable cannot be consistent over the training and test sets that I plan to train the models on. I will drop this variable

```
[152]: X.drop('last_pymnt_d', axis = 1, inplace = True)
```

next_pymnt_d: The next time that the borrower is expected to pay lending club in mm-yyyy

```
[153]: print(X['next_pymnt_d'].isna().sum())
X['next_pymnt_d'].head(10)
```

252971

```
[153]: 0    NaN
1    NaN
2    NaN
3    NaN
4    Feb-2016
5    NaN
6    Feb-2016
```

```

7      NaN
8      NaN
9      NaN
Name: next_pymnt_d, dtype: object

```

This variable has a large amount of missing values from all of the loans that have been fully paid or charged off. It is not a variable that will be consistent over the training and test sets.

```
[154]: X.drop('next_pymnt_d', axis = 1, inplace = True)
```

last_credit_pull_d: The last time lending club pulled the credit score of the borrower mm-yyyy

```
[155]: print(X['last_credit_pull_d'].isna().sum())
print(X['last_credit_pull_d'].dtype)
X['last_credit_pull_d'].head()
```

```

53
object

```

```
[155]: 0    Jan-2016
1    Sep-2013
2    Jan-2016
3    Jan-2015
4    Jan-2016
Name: last_credit_pull_d, dtype: object

```

There are 53 missing values that I can fill by the issue date because lending club pulls credit for every single loan so presumably there should have been a credit pull around the issue date.

```
[156]: # Fill in missing dates for the last credit pull column
X['last_credit_pull_d'] = X['last_credit_pull_d'].fillna(loan['issue_d'])
```

Similar to what I did with `earliest_cr_line` I will convert this column to numeric so that I can analyze its multicollinearity and feature importance.

```
[157]: # Convert the date that the lending company got a credit report
X['last_credit_pull_d'] = pd.to_datetime(X['last_credit_pull_d'])

X['last_credit_pull_d'] = X['last_credit_pull_d'].map(dt.datetime.toordinal)
```

application_type: Whether or not the loan application is joint or individual

```
[158]: X['application_type'].unique()
```

```
[158]: array(['INDIVIDUAL', 'JOINT'], dtype=object)
```

Since the only two values are 'INDIVIDUAL' and 'JOINT', I can use binary encoding to transform it into a numerical column by having 'INDIVIDUAL' -> 0 and 'JOINT' -> 1

```
[159]: # Convert pymnt_plan to a binary variable
X['application_type'] = np.where(X['application_type'] == 'JOINT', 1, 0)
```

verification_status_joint: Same as the verification_status but for joint applications.

```
[160]: X['verification_status_joint'].unique()
```

```
[160]: array([nan, 'Not Verified', 'Verified', 'Source Verified'], dtype=object)
```

The reason why there are missing values in this column is because most of the loans are to individual people. I can fill in the nan with some value that will allow me to use the pandas function `get_dummies`

```
[161]: # Fill the missing values in the application.
X['verification_status_joint'].fillna('indiv')

# One-hot encode the joint verification status
X = pd.get_dummies(X, columns = ['verification_status_joint'], drop_first =
↪False)
```

0.1.2 Columns to drop:

- **id:** This is not a relational database so there is no value in having a variable that can tie this dataset to others.
- **member_id:** same as ID, there are no other datasets for that need a bridge variable.

```
[162]: # drop all of the columns that I had mentioned earlier.
X.drop(['id', 'member_id'], axis = 1, inplace = True)
```

```
[163]: X.head()
```

```
[163]:
```

	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	\
0	5000.0	5000.0	4975.0	3	10.65	162.87	
1	2500.0	2500.0	2500.0	5	15.27	59.83	
2	2400.0	2400.0	2400.0	3	15.96	84.33	
3	10000.0	10000.0	10000.0	3	13.49	339.31	
4	3000.0	3000.0	3000.0	5	12.69	67.79	

	grade	emp_length	annual_inc	pymnt_plan	dti	delinq_2yrs	\
0	5.8	10.0	24000.0	0	27.65	0.0	
1	4.4	0.5	30000.0	0	1.00	0.0	
2	4.2	10.0	12252.0	0	8.72	0.0	
3	5.0	10.0	49200.0	0	20.00	0.0	
4	5.2	1.0	80000.0	0	17.94	0.0	

	earliest_cr_line	inq_last_6mths	mths_since_last_delinq	\
0	724642	1.0	NaN	
1	729845	5.0	NaN	
2	730790	2.0	NaN	

3	728690	1.0	35.0
4	728659	0.0	38.0

	mths_since_last_record	open_acc	pub_rec	revol_bal	revol_util \
0	NaN	3.0	0.0	13648.0	83.7
1	NaN	3.0	0.0	1687.0	9.4
2	NaN	2.0	0.0	2956.0	98.5
3	NaN	10.0	0.0	5598.0	21.0
4	NaN	15.0	0.0	27783.0	53.9

	total_acc	initial_list_status	last_pymnt_amnt	last_credit_pull_d \
0	9.0	0	171.62	735964
1	4.0	0	119.66	735112
2	10.0	0	649.91	735964
3	37.0	0	357.48	735599
4	38.0	0	67.79	735964

	collections_12_mths_ex_med	mths_since_last_major_derog	application_type \
0	0.0	NaN	0
1	0.0	NaN	0
2	0.0	NaN	0
3	0.0	NaN	0
4	0.0	NaN	0

	annual_inc_joint	dti_joint	acc_now_delinq	tot_coll_amt	tot_cur_bal \
0	NaN	NaN	0.0	NaN	NaN
1	NaN	NaN	0.0	NaN	NaN
2	NaN	NaN	0.0	NaN	NaN
3	NaN	NaN	0.0	NaN	NaN
4	NaN	NaN	0.0	NaN	NaN

	open_acc_6m	open_il_6m	open_il_12m	open_il_24m	mths_since_rcnt_il \
0	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN

	total_bal_il	il_util	open_rv_12m	open_rv_24m	max_bal_bc	all_util \
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN

	total_rev_hi_lim	inq_fi	total_cu_tl	inq_last_12m	emp_type_Accountant \
0	NaN	NaN	NaN	NaN	0

1	NaN	NaN	NaN	NaN	0
2	NaN	NaN	NaN	NaN	0
3	NaN	NaN	NaN	NaN	0
4	NaN	NaN	NaN	NaN	0

	emp_type_Admin	emp_type_Analyst	emp_type_Assistant	emp_type_Clergy	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	emp_type_Clerk	emp_type_Designer	emp_type_Director	emp_type_Education	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	emp_type_Executive	emp_type_Finance	emp_type_Healer	emp_type_Law	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	emp_type_Manager	emp_type_Manlab	emp_type_Operator	emp_type_Sales	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	emp_type_Service	emp_type_Technical	emp_type_Unemployed	emp_type_Vol	\
0	0	0	1	0	
1	0	0	0	0	
2	0	0	1	0	
3	0	0	0	0	
4	0	1	0	0	

	emp_type_other	home_ownership_OTHER	home_ownership_OWN	\
0	0	0	0	
1	1	0	0	
2	0	0	0	
3	1	0	0	
4	0	0	0	

	home_ownership_RENT	verification_status_Source Verified	\
0	1	0	
1	1	1	
2	1	0	
3	1	1	
4	1	1	

	verification_status_Verified	purpose_credit_card	\
0	1	1	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	purpose_debt_consolidation	purpose_educational	purpose_home_improvement	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	purpose_house	purpose_major_purchase	purpose_medical	purpose_moving	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	purpose_other	purpose_renewable_energy	purpose_small_business	\
0	0	0	0	
1	0	0	0	
2	0	0	1	
3	1	0	0	
4	1	0	0	

	purpose_vacation	purpose_wedding	addr_state_AL	addr_state_AR	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	addr_state_AZ	addr_state_CA	addr_state_CO	addr_state_CT	addr_state_DC	\
0	1	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	1	0	0	0	

4	0	0	0	0	0
	addr_state_DE	addr_state_FL	addr_state_GA	addr_state_HI	addr_state_IA \
0	0	0	0	0	0
1	0	0	1	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
	addr_state_ID	addr_state_IL	addr_state_IN	addr_state_KS	addr_state_KY \
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
	addr_state_LA	addr_state_MA	addr_state_MD	addr_state_ME	addr_state_MI \
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
	addr_state_MN	addr_state_MO	addr_state_MS	addr_state_MT	addr_state_NC \
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
	addr_state_ND	addr_state_NE	addr_state_NH	addr_state_NJ	addr_state_NM \
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
	addr_state_NV	addr_state_NY	addr_state_OH	addr_state_OK	addr_state_OR \
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	1
	addr_state_PA	addr_state_RI	addr_state_SC	addr_state_SD	addr_state_TN \
0	0	0	0	0	0
1	0	0	0	0	0

2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

	addr_state_TX	addr_state_UT	addr_state_VA	addr_state_VT	addr_state_WA	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	addr_state_WI	addr_state_WV	addr_state_WY	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	verification_status_joint_Not Verified	\
0	0	
1	0	
2	0	
3	0	
4	0	


	verification_status_joint_Source Verified	\
0	0	
1	0	
2	0	
3	0	
4	0	

	verification_status_joint_Verified
0	0
1	0
2	0
3	0
4	0

0.2 Missing Values

```
[164]: plt.figure(figsize = (30,7));
# Plot missing values for each column
plt.bar(x = X.columns, height = X.isna().sum()/X.shape[0])
# hide the names of the columns for a cleaner graph
plt.xticks([]);
# Set the scale to log so I can see every column that has a missing value
```

```
plt.yscale('log');
```



```
[165]: list(X.columns[X.isna().sum()/X.shape[0] > 0])
```

```
[165]: ['annual_inc',  
        'delinq_2yrs',  
        'inq_last_6mths',  
        'mths_since_last_delinq',  
        'mths_since_last_record',  
        'open_acc',  
        'pub_rec',  
        'revol_util',  
        'total_acc',  
        'collections_12_mths_ex_med',  
        'mths_since_last_major_derog',  
        'annual_inc_joint',  
        'dti_joint',  
        'acc_now_delinq',  
        'tot_coll_amt',  
        'tot_cur_bal',  
        'open_acc_6m',  
        'open_il_6m',  
        'open_il_12m',  
        'open_il_24m',  
        'mths_since_rcnt_il',  
        'total_bal_il',  
        'il_util',  
        'open_rv_12m',  
        'open_rv_24m',  
        'max_bal_bc',  
        'all_util',  
        'total_rev_hi_lim',  
        'inq_fi',  
        'total_cu_tl',  
        'inq_last_12m']
```

The imputation of values for this dataset is the most challenging aspect of this project. There are

many columns that have large swaths of missing values that need reasonable values filled in so as to allow the algorithms to capture the existing patterns in the data instead of artificial ones coming from synthetic data.

My assumption was that most missing values were zeroes that people forgot to fill out on forms or didn't bother with. For instance 'open_il_12m' refers to the number of opened installment accounts in the last 12 months. If there isn't any information on the number of installment accounts, then there likely aren't any installment accounts.

Values such as 'mths_since_last_delinq' or months since last delinquent are different because if someone was never delinquent then the correct answer isn't 0 which would imply that they are currently delinquent but a very high number. The higher the value, the better the outcome on their loan. I can use 1000 since it is 83 years worth of non-delinquency.

```
[166]: # Create a list with all of the columns that have missing values to be filled
        ↪with 0
fill_with_zero = ['annual_inc', 'delinq_2yrs', 'inq_last_6mths', 'open_acc',
                  'pub_rec', 'revol_util', 'total_acc',
                  ↪'collections_12_mths_ex_med',
                  'acc_now_delinq', 'tot_coll_amt', 'tot_cur_bal',
                  ↪'open_acc_6m',
                  'open_il_6m', 'open_il_12m', 'open_il_24m', 'total_bal_il',
                  ↪'il_util',
                  'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util',
                  ↪'total_rev_hi_lim',
                  'inq-fi', 'total_cu_tl', 'inq_last_12m']

# Fill those columns with 0
X[fill_with_zero] = X[fill_with_zero].fillna(0)
```

```
[167]: # Create a list with all of the columns that have missing values to be filled
        ↪with 1000
fill_with_1000 = ['mths_since_last_delinq', 'mths_since_last_record',
                  'mths_since_last_major_derog', 'mths_since_rcnt_il']

# Fill those columns with 0
X[fill_with_1000] = X[fill_with_1000].fillna(1000)
```

```
[168]: # List out all the columns that still having missing values
list(X.columns[X.isna().sum()/X.shape[0] > 0])
```

```
[168]: ['annual_inc_joint', 'dti_joint']
```

```
[169]: # Combine annual inc columns joint and indiv into one.
X['annual_inc'] = [loan['annual_inc'][i] if math.isnan(x) else x for i, x in
                  ↪enumerate(loan['annual_inc_joint'])]

# Create a column that documents if the column above is joint or indiv
X['inc_joint'] = np.where(loan['annual_inc_joint'].isnull(), 0, 1)
```

```
[170]: # Combine dti columns joint and indiv into one.
X['dti'] = [X['dti'][i] if math.isnan(x) else x for i, x in enumerate(X['dti_joint'])]
# Create a column that documents if the column above is joint or indiv
X['dti_joint'] = np.where(loan['dti_joint'].isnull(), 0, 1)

[171]: # Check to see if there are still missing columns
list(X.columns[X.isna().sum()/X.shape[0] > 0])

[171]: ['annual_inc', 'annual_inc_joint']

[172]: # Drop 'annual_inc_joint' now that it has been added to 'annual_inc'
X.drop('annual_inc_joint', axis = 1, inplace = True)
# Fill in the missing values with 0 as the assumption is there is no income
X['annual_inc'].fillna(0, inplace = True)

[173]: # Check to see if there are still missing columns
list(X.columns[X.isna().sum()/X.shape[0] > 0])

[173]: []

[174]: # Initialize the standard scaler
scaler = StandardScaler()
# Fit and transform X into a dataframe so that I can keep using pandas
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns = X.columns)

[175]: X_scaled.head()
```

	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	\
0	-1.156460	-1.155635	-1.152256	-0.654724	-0.592611	-1.121467	
1	-1.452829	-1.452198	-1.445430	1.527360	0.461735	-1.543440	
2	-1.464683	-1.464061	-1.457275	-0.654724	0.619202	-1.443107	
3	-0.563724	-0.562507	-0.557025	-0.654724	0.055515	-0.398905	
4	-1.393555	-1.392886	-1.386203	1.527360	-0.127055	-1.510842	

	grade	emp_length	annual_inc	pymnt_plan	dti	delinq_2yrs	\
0	0.763889	1.134870	-0.789014	-0.003357	1.147690	-0.364672	
1	-0.314193	-1.398891	-0.696292	-0.003357	-2.063920	-0.364672	
2	-0.468204	1.134870	-0.970563	-0.003357	-1.133577	-0.364672	
3	0.147842	1.134870	-0.399582	-0.003357	0.225784	-0.364672	
4	0.301854	-1.265535	0.076390	-0.003357	-0.022468	-0.364672	

	earliest_cr_line	inq_last_6mths	mths_since_last_delinq	\
0	-1.771866	0.305877	0.975849	
1	0.143584	4.312132	0.975849	
2	0.491479	1.307441	0.975849	
3	-0.281622	0.305877	-1.021784	

4	-0.293034	-0.695687	-1.015574
---	-----------	-----------	-----------

	mths_since_last_record	open_acc	pub_rec	revol_bal	revol_util	\
0	0.427154	-1.607499	-0.335522	-0.145932	1.201140	
1	0.427154	-1.607499	-0.335522	-0.679268	-1.912396	
2	0.427154	-1.795553	-0.335522	-0.622684	1.821333	
3	0.427154	-0.291124	-0.335522	-0.504878	-1.426299	
4	0.427154	0.649144	-0.335522	0.484341	-0.047627	

	total_acc	initial_list_status	last_pymnt_amnt	last_credit_pull_d	\
0	-1.373775	-0.97077	-0.415561	0.302329	
1	-1.796028	-0.97077	-0.426398	-3.540893	
2	-1.289324	-0.97077	-0.315809	0.302329	
3	0.990842	-0.97077	-0.376798	-1.344122	
4	1.075293	-0.97077	-0.437216	0.302329	

	collections_12_mths_ex_med	mths_since_last_major_derog	application_type	\
0	-0.107149		0.576897	-0.024004
1	-0.107149		0.576897	-0.024004
2	-0.107149		0.576897	-0.024004
3	-0.107149		0.576897	-0.024004
4	-0.107149		0.576897	-0.024004

	dti_joint	acc_now_delinq	tot_coll_amt	tot_cur_bal	open_acc_6m	\
0	-0.023957	-0.064298	-0.021004	-0.843347	-0.103891	
1	-0.023957	-0.064298	-0.021004	-0.843347	-0.103891	
2	-0.023957	-0.064298	-0.021004	-0.843347	-0.103891	
3	-0.023957	-0.064298	-0.021004	-0.843347	-0.103891	
4	-0.023957	-0.064298	-0.021004	-0.843347	-0.103891	

	open_il_6m	open_il_12m	open_il_24m	mths_since_rcnt_il	total_bal_il	\
0	-0.107375	-0.094692	-0.109934	0.154904	-0.100884	
1	-0.107375	-0.094692	-0.109934	0.154904	-0.100884	
2	-0.107375	-0.094692	-0.109934	0.154904	-0.100884	
3	-0.107375	-0.094692	-0.109934	0.154904	-0.100884	
4	-0.107375	-0.094692	-0.109934	0.154904	-0.100884	

	il_util	open_rv_12m	open_rv_24m	max_bal_bc	all_util	total_rev_hi_lim	\
0	-0.139209	-0.105267	-0.117038	-0.116273	-0.149047	-0.79786	
1	-0.139209	-0.105267	-0.117038	-0.116273	-0.149047	-0.79786	
2	-0.139209	-0.105267	-0.117038	-0.116273	-0.149047	-0.79786	
3	-0.139209	-0.105267	-0.117038	-0.116273	-0.149047	-0.79786	
4	-0.139209	-0.105267	-0.117038	-0.116273	-0.149047	-0.79786	

	inq-fi	total_cu_tl	inq_last_12m	emp_type_Accountant	emp_type_Admin	\
0	-0.085105	-0.076504	-0.088306	-0.109205	-0.104843	
1	-0.085105	-0.076504	-0.088306	-0.109205	-0.104843	

2	-0.085105	-0.076504	-0.088306	-0.109205	-0.104843
3	-0.085105	-0.076504	-0.088306	-0.109205	-0.104843
4	-0.085105	-0.076504	-0.088306	-0.109205	-0.104843

	emp_type_Analyst	emp_type_Assistant	emp_type_Clergy	emp_type_Clerk	\
0	-0.157505	-0.147428	-0.039637	-0.077757	
1	-0.157505	-0.147428	-0.039637	-0.077757	
2	-0.157505	-0.147428	-0.039637	-0.077757	
3	-0.157505	-0.147428	-0.039637	-0.077757	
4	-0.157505	-0.147428	-0.039637	-0.077757	

	emp_type_Designer	emp_type_Director	emp_type_Education	\
0	-0.052789	-0.196448	-0.186479	
1	-0.052789	-0.196448	-0.186479	
2	-0.052789	-0.196448	-0.186479	
3	-0.052789	-0.196448	-0.186479	
4	-0.052789	-0.196448	-0.186479	

	emp_type_Executive	emp_type_Finance	emp_type_Healer	emp_type_Law	\
0	-0.180873	-0.117117	-0.245189	-0.046832	
1	-0.180873	-0.117117	-0.245189	-0.046832	
2	-0.180873	-0.117117	-0.245189	-0.046832	
3	-0.180873	-0.117117	-0.245189	-0.046832	
4	-0.180873	-0.117117	-0.245189	-0.046832	

	emp_type_Manager	emp_type_Manlab	emp_type_Operator	emp_type_Sales	\
0	-0.435078	-0.082293	-0.169413	-0.142468	
1	-0.435078	-0.082293	-0.169413	-0.142468	
2	-0.435078	-0.082293	-0.169413	-0.142468	
3	-0.435078	-0.082293	-0.169413	-0.142468	
4	-0.435078	-0.082293	-0.169413	-0.142468	

	emp_type_Service	emp_type_Technical	emp_type_Unemployed	emp_type_Vol	\
0	-0.160378	-0.395006	4.030308	-0.172517	
1	-0.160378	-0.395006	-0.248120	-0.172517	
2	-0.160378	-0.395006	4.030308	-0.172517	
3	-0.160378	-0.395006	-0.248120	-0.172517	
4	-0.160378	2.531605	-0.248120	-0.172517	

	emp_type_other	home_ownership_OTHER	home_ownership_OWN	\
0	-0.631040	-0.016276	-0.330681	
1	1.584686	-0.016276	-0.330681	
2	-0.631040	-0.016276	-0.330681	
3	1.584686	-0.016276	-0.330681	
4	-0.631040	-0.016276	-0.330681	

home_ownership_RENT	verification_status_Source	Verified	\
---------------------	----------------------------	----------	---

0	1.2214	-0.768632
1	1.2214	1.301013
2	1.2214	-0.768632
3	1.2214	1.301013
4	1.2214	1.301013

	verification_status_Verified	purpose_credit_card \
0	1.431317	1.817653
1	-0.698657	-0.550160
2	-0.698657	-0.550160
3	-0.698657	-0.550160
4	-0.698657	-0.550160

	purpose_debt_consolidation	purpose_educational	purpose_home_improvement \
0	-1.201443	-0.021838	-0.249058
1	-1.201443	-0.021838	-0.249058
2	-1.201443	-0.021838	-0.249058
3	-1.201443	-0.021838	-0.249058
4	-1.201443	-0.021838	-0.249058

	purpose_house	purpose_major_purchase	purpose_medical	purpose_moving \
0	-0.064769	-0.140912	-0.098577	-0.078349
1	-0.064769	-0.140912	-0.098577	-0.078349
2	-0.064769	-0.140912	-0.098577	-0.078349
3	-0.064769	-0.140912	-0.098577	-0.078349
4	-0.064769	-0.140912	-0.098577	-0.078349

	purpose_other	purpose_renewable_energy	purpose_small_business \
0	-0.225373	-0.025464	-0.108777
1	-0.225373	-0.025464	-0.108777
2	-0.225373	-0.025464	9.193151
3	4.437084	-0.025464	-0.108777
4	4.437084	-0.025464	-0.108777

	purpose_vacation	purpose_wedding	addr_state_AL	addr_state_AR \
0	-0.073251	-0.051496	-0.113061	-0.086828
1	-0.073251	-0.051496	-0.113061	-0.086828
2	-0.073251	-0.051496	-0.113061	-0.086828
3	-0.073251	-0.051496	-0.113061	-0.086828
4	-0.073251	-0.051496	-0.113061	-0.086828

	addr_state_AZ	addr_state_CA	addr_state_CO	addr_state_CT	addr_state_DC \
0	6.517162	-0.413398	-0.147149	-0.124436	-0.052423
1	-0.153441	-0.413398	-0.147149	-0.124436	-0.052423
2	-0.153441	-0.413398	-0.147149	-0.124436	-0.052423
3	-0.153441	2.418977	-0.147149	-0.124436	-0.052423
4	-0.153441	-0.413398	-0.147149	-0.124436	-0.052423

	addr_state_DE	addr_state_FL	addr_state_GA	addr_state_HI	addr_state_IA	\
0	-0.05327	-0.271536	-0.184084	-0.071949	-0.003972	
1	-0.05327	-0.271536	5.432297	-0.071949	-0.003972	
2	-0.05327	-0.271536	-0.184084	-0.071949	-0.003972	
3	-0.05327	-0.271536	-0.184084	-0.071949	-0.003972	
4	-0.05327	-0.271536	-0.184084	-0.071949	-0.003972	

	addr_state_ID	addr_state_IL	addr_state_IN	addr_state_KS	addr_state_KY	\
0	-0.003677	-0.204067	-0.125636	-0.094934	-0.098635	
1	-0.003677	-0.204067	-0.125636	-0.094934	-0.098635	
2	-0.003677	4.900357	-0.125636	-0.094934	-0.098635	
3	-0.003677	-0.204067	-0.125636	-0.094934	-0.098635	
4	-0.003677	-0.204067	-0.125636	-0.094934	-0.098635	

	addr_state_LA	addr_state_MA	addr_state_MD	addr_state_ME	addr_state_MI	\
0	-0.109885	-0.154136	-0.155806	-0.024331	-0.163067	
1	-0.109885	-0.154136	-0.155806	-0.024331	-0.163067	
2	-0.109885	-0.154136	-0.155806	-0.024331	-0.163067	
3	-0.109885	-0.154136	-0.155806	-0.024331	-0.163067	
4	-0.109885	-0.154136	-0.155806	-0.024331	-0.163067	

	addr_state_MN	addr_state_MO	addr_state_MS	addr_state_MT	addr_state_NC	\
0	-0.13532	-0.127556	-0.065744	-0.053768	-0.16928	
1	-0.13532	-0.127556	-0.065744	-0.053768	-0.16928	
2	-0.13532	-0.127556	-0.065744	-0.053768	-0.16928	
3	-0.13532	-0.127556	-0.065744	-0.053768	-0.16928	
4	-0.13532	-0.127556	-0.065744	-0.053768	-0.16928	

	addr_state_ND	addr_state_NE	addr_state_NH	addr_state_NJ	addr_state_NM	\
0	-0.02324	-0.036428	-0.069732	-0.197322	-0.074813	
1	-0.02324	-0.036428	-0.069732	-0.197322	-0.074813	
2	-0.02324	-0.036428	-0.069732	-0.197322	-0.074813	
3	-0.02324	-0.036428	-0.069732	-0.197322	-0.074813	
4	-0.02324	-0.036428	-0.069732	-0.197322	-0.074813	

	addr_state_NV	addr_state_NY	addr_state_OH	addr_state_OK	addr_state_OR	\
0	-0.119254	-0.301818	-0.185863	-0.09589	-0.111481	
1	-0.119254	-0.301818	-0.185863	-0.09589	-0.111481	
2	-0.119254	-0.301818	-0.185863	-0.09589	-0.111481	
3	-0.119254	-0.301818	-0.185863	-0.09589	-0.111481	
4	-0.119254	-0.301818	-0.185863	-0.09589	8.970130	

	addr_state_PA	addr_state_RI	addr_state_SC	addr_state_SD	addr_state_TN	\
0	-0.191506	-0.066381	-0.110158	-0.045272	-0.121394	
1	-0.191506	-0.066381	-0.110158	-0.045272	-0.121394	
2	-0.191506	-0.066381	-0.110158	-0.045272	-0.121394	

3	-0.191506	-0.066381	-0.110158	-0.045272	-0.121394
4	-0.191506	-0.066381	-0.110158	-0.045272	-0.121394

	addr_state_TX	addr_state_UT	addr_state_VA	addr_state_VT	addr_state_WA \
0	-0.295217	-0.084316	-0.174612	-0.045046	-0.149636
1	-0.295217	-0.084316	-0.174612	-0.045046	-0.149636
2	-0.295217	-0.084316	-0.174612	-0.045046	-0.149636
3	-0.295217	-0.084316	-0.174612	-0.045046	-0.149636
4	-0.295217	-0.084316	-0.174612	-0.045046	-0.149636

	addr_state_WI	addr_state_WV	addr_state_WY \
0	-0.114958	-0.070478	-0.04786
1	-0.114958	-0.070478	-0.04786
2	-0.114958	-0.070478	-0.04786
3	-0.114958	-0.070478	-0.04786
4	-0.114958	-0.070478	-0.04786

	verification_status_joint_Not Verified \
0	-0.017861
1	-0.017861
2	-0.017861
3	-0.017861
4	-0.017861

	verification_status_joint_Source Verified \
0	-0.008291
1	-0.008291
2	-0.008291
3	-0.008291
4	-0.008291

	verification_status_joint_Verified	inc_joint
0	-0.01372	-0.024004
1	-0.01372	-0.024004
2	-0.01372	-0.024004
3	-0.01372	-0.024004
4	-0.01372	-0.024004

Multicollinearity

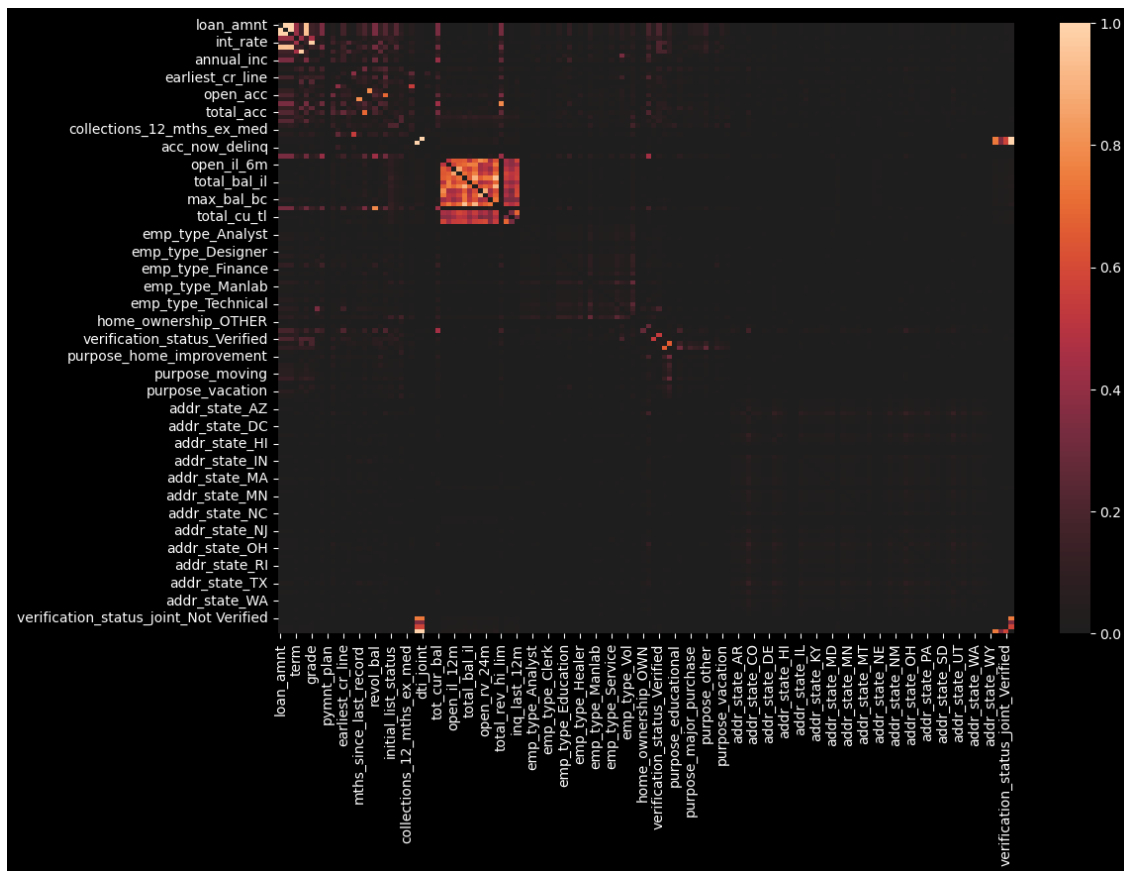
```
[176]: # Function to get rid of all the diagonals in correlation heatmaps
def zero_diagonal_heatmap(dfx, min = 0, max = None, color_map = None,
↪ threshold = 0):
    # give the max a base value
    if max is None:
        max = dfx.shape[1]
    # Create a numpy matrix that is easy to manipulate
    cm = np.array(abs(dfx.iloc[:,min:max].corr(numeric_only = False)))
```

```


# Create a loop that replaces all the diagonal values with 0
for i in range(len(cm)):
    # Loop thorough the other axis of values
    for j in range(len(cm)):
        # All diagonal values happen when i == j
        if i == j:
            # set the diagonal value equal to 0
            cm[i,j] = 0
    df_cm = pd.DataFrame(cm, columns = list(dfx.columns)[min:max], index =
    ↪list(dfx.columns)[min:max])
    # Make the plot bigger
    plt.figure(figsize=(12, 8))
    # Plot the heatmap
    ax = sns.heatmap(df_cm, cmap = color_map, center = threshold, vmin = 0,
    ↪vmax = 1);

```

```
[177]: zero_diagonal_heatmap(X_scaled)
```



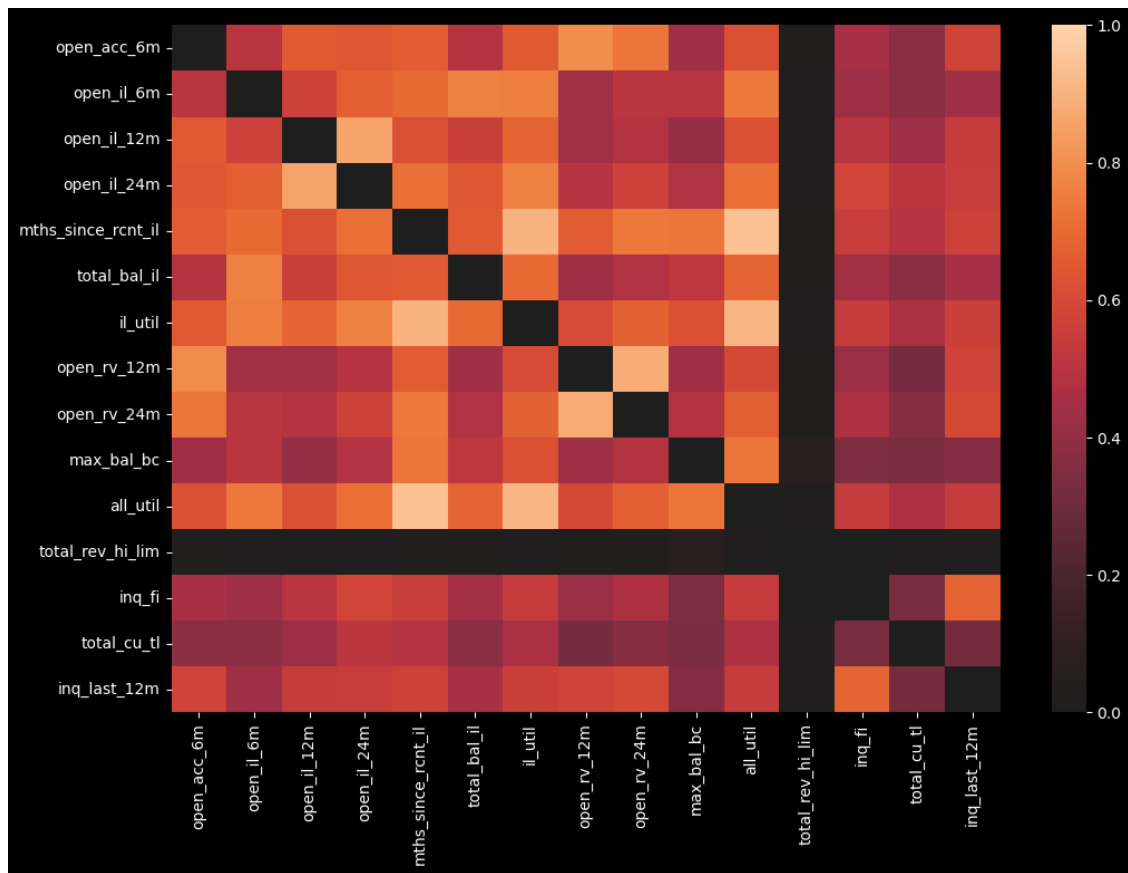
```
[178]: # Create a heatmap of the first 7 columns
zero_diagonal_heatmap(X_scaled, 0, 7)
```



'loan_amnt' is the amount the requested amount of money for the loan from the borrower, 'funded_amnt' is the amount of money that has been committed to the loan, and 'funded_amnt_inv' is the amount of money investors have committed to the loan. Since loan amount is the only variable that the applicant has control over, I believe it is the only relevant variable to keep. 'installment' is the monthly payment if the loan originates, this is a function of the loan amount and the term and can be dropped. 'int_rate' is the interest rate of the loan and it is directly determined by the grade that lending club gives the loan. I can pick one of the two and it will contain the information of the other.

```
[179]: # Drop all the columns that are clearly collinear that are the least interesting
X_scaled.drop(['funded_amnt', 'funded_amnt_inv', 'installment', 'int_rate'], axis=
↳ 1, inplace = True)
```

```
[180]: zero_diagonal_heatmap(X_scaled, 27, 42)
```



```
[181]: # Create a function that takes a dataframe and a correlation level
def corr_dict(dfx, corr = 0.7):
    # Create the correlation with absolute values to measure collinearity
    corr_mat = abs(dfx.corr(numeric_only = False)) > corr # type: ignore
    # Define the dictionary outside the function for uses
    global correl_dict
    # Create a dictionary with sorted values so that it is easy to see the next
    ↪ steps
    correl_dict = dict(corr_mat.sum().sort_values(ascending = False))
    # Subtract each value by one for the self correlation on the diagonal
    correl_dict = {key: value - 1 for key, value in correl_dict.items() if
    ↪ value != 1}
    # Print the dictionary
    print(correl_dict)
```

```
[182]: # Create a function that takes a list of features and minimum number of
    ↪ correlations
def corr_list(min_corr, corr, still_high_mcl = None):
    # Define the dictionary where all the information will be stored
    correlations2 = {}
```

```

# Almost always the default value will be None
if still_high_mcl is None:
    # # the default list is the keys from correl_dict whose values are at_
    ↪ least min_corr
    still_high_mcl = [key for key, value in correl_dict.items() if value >=
    ↪ min_corr]
    # Drop all date time objects from the list by using the ending _d as an_
    ↪ identifier
    still_high_mcl = [element for element in still_high_mcl if not element.
    ↪ endswith("_d")]
    # Store the correlation matrix in a variable for easy use
    still_corr = abs(X_scaled[still_high_mcl].corr(numeric_only = True))
    # Put the index and column names as the feature names
    still_corr.columns = still_high_mcl
    still_corr.index = still_high_mcl
    # Iterate through all the features in the list
    for shm in still_high_mcl:
        # create a list of all the features that have a correlation of over 7
        corrs2 = still_corr.index[still_corr[shm] > corr].tolist() #type: ignore
        # remove the self correlation
        corrs2.remove(shm)
        # input the feature and list into the dictionary
        correlations2.update({shm: corrs2})
    # print the result
    print(correlations2)

```

```

[183]: # Find the number of features each feature is correlated with
corr_dict(X_scaled)

```

```

{'all_util': 5, 'il_util': 5, 'mths_since_rcnt_il': 5, 'open_il_24m': 4,
'inc_joint': 3, 'application_type': 3, 'verification_status_joint_Not Verified':
3, 'open_rv_24m': 3, 'open_il_6m': 3, 'dti_joint': 3, 'max_bal_bc': 2,
'open_rv_12m': 2, 'total_bal_il': 2, 'open_acc_6m': 2, 'open_il_12m': 1,
'total_rev_hi_lim': 1, 'mths_since_last_record': 1, 'pub_rec': 1, 'revol_bal':
1}

```

```

[184]: # Get the list of features each feature is correlated with.
corr_list(corr = 0.7, min_corr = 1)

```

```

{'all_util': ['il_util', 'mths_since_rcnt_il', 'open_il_24m', 'open_il_6m',
'max_bal_bc'], 'il_util': ['all_util', 'mths_since_rcnt_il', 'open_il_24m',
'open_il_6m', 'total_bal_il'], 'mths_since_rcnt_il': ['all_util', 'il_util',
'open_il_24m', 'open_rv_24m', 'max_bal_bc'], 'open_il_24m': ['all_util',
'il_util', 'mths_since_rcnt_il', 'open_il_12m'], 'inc_joint':
['application_type', 'verification_status_joint_Not Verified', 'dti_joint'],
'application_type': ['inc_joint', 'verification_status_joint_Not Verified',
'dti_joint'], 'verification_status_joint_Not Verified': ['inc_joint',
'application_type', 'dti_joint'], 'open_rv_24m': ['mths_since_rcnt_il',

```

```
'open_rv_12m', 'open_acc_6m'], 'open_il_6m': ['all_util', 'il_util',
'total_bal_il'], 'dti_joint': ['inc_joint', 'application_type',
'verification_status_joint_Not Verified'], 'max_bal_bc': ['all_util',
'mths_since_rcnt_il'], 'open_rv_12m': ['open_rv_24m', 'open_acc_6m'],
'total_bal_il': ['il_util', 'open_il_6m'], 'open_acc_6m': ['open_rv_24m',
'open_rv_12m'], 'open_il_12m': ['open_il_24m'], 'total_rev_hi_lim':
['revol_bal'], 'mths_since_last_record': ['pub_rec'], 'pub_rec':
['mths_since_last_record'], 'revol_bal': ['total_rev_hi_lim']}]
```

One set of features that is consistently multicollinear is the features that are measured in 6 months, 12 months, and 24 months. It isn't surprising that they are collinear since they are measuring the same thing but it is important to combine them so that I can preserve their information without confusing their impact. The simplest way to do this is to set the time frame equal and then compare.

```
[185]: # Scale the number of 6 month installment accounts to 12
X_scaled['il_6'] = X_scaled['open_il_6m'] * 2
# Scale the number of 24 month installment accounts to 12
X_scaled['il_24'] = X_scaled['open_il_24m'] / 2
# Take the max out of all the scaled features to find if there is a spike in
↳install accounts.
X_scaled['open_install_acc'] = X_scaled[['il_6', 'open_il_12m', 'il_24']].
↳max(axis = 1)

# Scale the number of 24 month revolving accounts to 12
X_scaled['rv_24'] = X_scaled['open_rv_24m'] / 2
# Take the max out of all the scaled features to find if there is a spike in
↳revolving accounts.
X_scaled['open_revolving'] = X_scaled[['open_rv_12m', 'rv_24']].max(axis = 1)

# Drop all the original columns now that they are useless.
X_scaled.drop(['il_6', 'open_il_12m', 'il_24', 'open_rv_12m', 'rv_24',
↳'open_il_6m', 'open_il_24m', 'open_rv_24m'], axis = 1, inplace = True)
```

```
[186]: # Find the number of features each feature is correlated with again.
corr_dict(X_scaled)

{'mths_since_rcnt_il': 5, 'open_install_acc': 4, 'il_util': 4, 'all_util': 4,
'inc_joint': 3, 'dti_joint': 3, 'application_type': 3,
'verification_status_joint_Not Verified': 3, 'open_revolving': 2,
'total_bal_il': 2, 'max_bal_bc': 2, 'revol_bal': 1, 'total_rev_hi_lim': 1,
'open_acc_6m': 1, 'pub_rec': 1, 'mths_since_last_record': 1}
```

```
[187]: # Get the list of features each feature is correlated with again.
corr_list(corr = 0.7, min_corr = 1)

{'mths_since_rcnt_il': ['open_install_acc', 'il_util', 'all_util',
'open_revolving', 'max_bal_bc'], 'open_install_acc': ['mths_since_rcnt_il',
'il_util', 'all_util', 'total_bal_il'], 'il_util': ['mths_since_rcnt_il',
'open_install_acc', 'all_util', 'total_bal_il'], 'all_util':
```

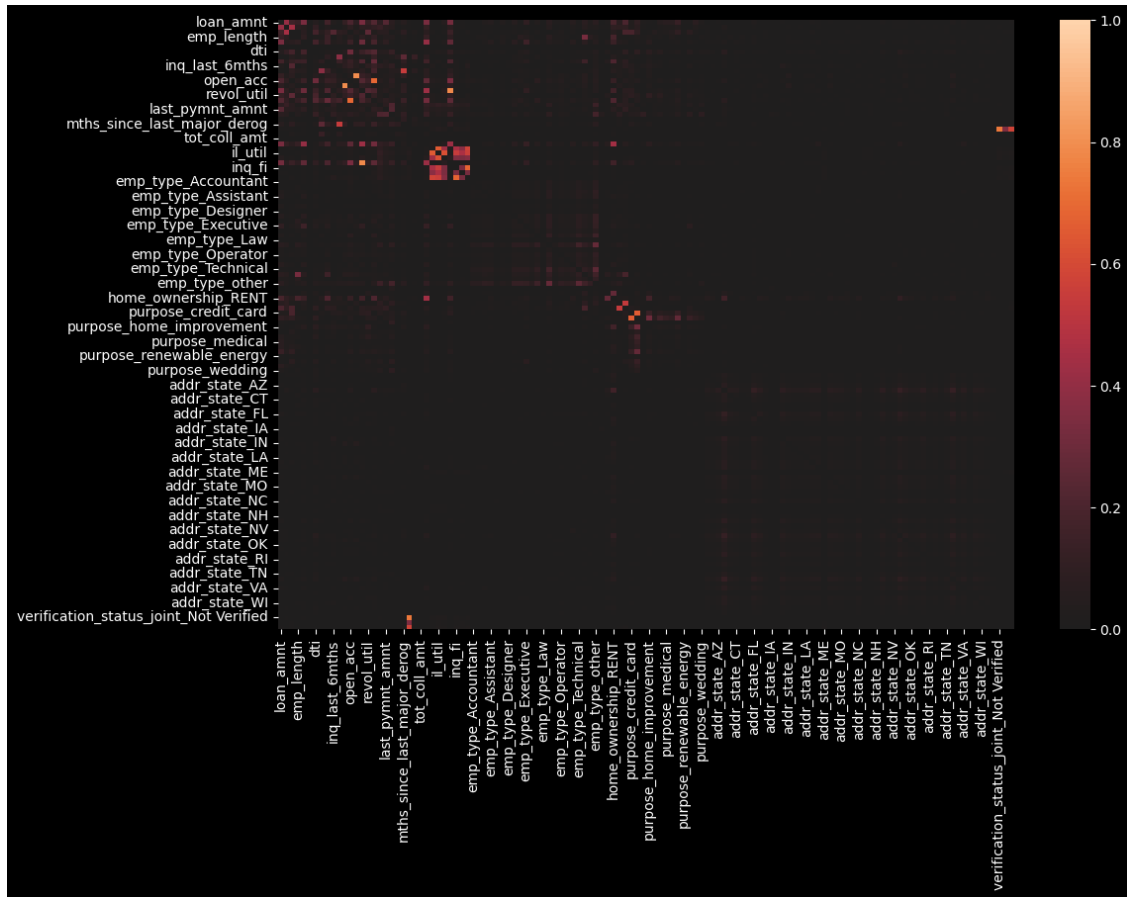
```
['mths_since_rcnt_il', 'open_install_acc', 'il_util', 'max_bal_bc'],
'inc_joint': ['dti_joint', 'application_type', 'verification_status_joint_Not
Verified'], 'dti_joint': ['inc_joint', 'application_type',
'verification_status_joint_Not Verified'], 'application_type': ['inc_joint',
'dti_joint', 'verification_status_joint_Not Verified'],
'verification_status_joint_Not Verified': ['inc_joint', 'dti_joint',
'application_type'], 'open_revolving': ['mths_since_rcnt_il', 'open_acc_6m'],
'total_bal_il': ['open_install_acc', 'il_util'], 'max_bal_bc':
['mths_since_rcnt_il', 'all_util'], 'revol_bal': ['total_rev_hi_lim'],
'total_rev_hi_lim': ['revol_bal'], 'open_acc_6m': ['open_revolving'], 'pub_rec':
['mths_since_last_record'], 'mths_since_last_record': ['pub_rec']}]
```

Criteria for choosing collinear features - features that have default values of 0 are better than features with undefined values. - dollar amounts over number of accounts. They give more information and once they are properly scaled they will be more useful. - The more information the variable contains the better. - Synthesized features are better than original features.

il_util max_bal_bc total_rev_hi_lim

```
[188]: # Drop the features that least fit the criteria
X_scaled.drop(['mths_since_rcnt_il', 'all_util', 'open_install_acc',
↳ 'open_revolving',
        'total_bal_il', 'inc_joint', 'application_type'], axis = 1,
↳ inplace = True)
```

```
[189]: zero_diagonal_heatmap(X_scaled)
```



```
[190]: corr_dict(X_scaled)

{'mths_since_last_record': 1, 'dti_joint': 1, 'verification_status_joint_Not
Verified': 1, 'revol_bal': 1, 'total_rev_hi_lim': 1, 'pub_rec': 1}
```

```
[191]: corr_list(corr = 0.7, min_corr = 1)

{'mths_since_last_record': ['pub_rec'], 'dti_joint':
['verification_status_joint_Not Verified'], 'verification_status_joint_Not
Verified': ['dti_joint'], 'revol_bal': ['total_rev_hi_lim'], 'total_rev_hi_lim':
['revol_bal'], 'pub_rec': ['mths_since_last_record']}
```

Unsurprisingly two features that both rely on the borrowers having a joint account are collinear. I will drop the dummy variable as the other one contains more information

```
[192]: # Drop the least relevant columns that are in corr_list
X_scaled.drop(['mths_since_last_record', 'total_rev_hi_lim',
               'verification_status_joint_Not Verified'], axis = 1, inplace = True)
```

To make sure that I had fully eliminated the multicollinearity I calculated every remaining feature's

Variance Inflation Factor (VIF). The formula for VIF is:

$$VIF = \frac{1}{1 - R^2}$$

Where R^2 is the R-squared value that represents correlations.

```
[193]: # Create vif function with a dataframe as the argument
def vif(dfx):
    #calculate the vif for the dataframe to see if it is non collinear
    vifs = [variance_inflation_factor(dfx.values, i) for i in range(dfx.
↪shape[1])]
    # Put those numbers into a dictionary
    vif_dict = {dfx.columns[i]:vifs[i] for i in range(dfx.shape[1])}
    # Create a filtered dictionary to find out how many features have a VIF
↪under 2.5
    filtered_dict = {k: v for k, v in vif_dict.items() if v < 2.5}
    # Print the amount of columns that aren't collinear out of all the columns
    print(f"{len(filtered_dict)} out of {len(dfx.columns)} features have a vif
↪< 2.5")
    # Create a sorted dictionary based on the values
    vif_dict = {k:v for k, v in sorted(vif_dict.items(), key=lambda item:
↪item[1], reverse = True)}
    # Print the result
    print(vif_dict)
```

```
[194]: vif(X_scaled)
```

```
57 out of 123 features have a vif < 2.5
{'emp_type_other': 169.57974002764598, 'emp_type_Manager': 111.7147586447447,
'emp_type_Technical': 97.59897197456667, 'addr_state_CA': 51.094645200048134,
'emp_type_Unemployed': 46.39212956418784, 'emp_type_Healer': 45.254717422269884,
'addr_state_NY': 31.767745462942365, 'addr_state_TX': 30.633566724669535,
'emp_type_Director': 30.59330240518125, 'emp_type_Education': 27.89717101959111,
'addr_state_FL': 26.705856680717247, 'emp_type_Executive': 26.408281585647327,
'purpose_debt_consolidation': 24.96764806132425, 'emp_type_Vol':
24.23857849854677, 'emp_type_Operator': 23.461031656313015, 'emp_type_Service':
21.245410815327098, 'emp_type_Analyst': 20.537956437604667,
'purpose_credit_card': 18.93967828897951, 'emp_type_Assistant':
18.239376685202956, 'emp_type_Sales': 17.13106743520082, 'addr_state_IL':
16.424458850979022, 'addr_state_NJ': 15.502262195704335, 'addr_state_PA':
14.716247800285, 'addr_state_OH': 13.980629680318021, 'addr_state_GA':
13.74480709611746, 'addr_state_VA': 12.536294915046794, 'emp_type_Finance':
12.043774061171954, 'addr_state_NC': 11.88644182350205, 'addr_state_MI':
11.147880276767328, 'emp_type_Accountant': 10.636077148271605, 'addr_state_MD':
10.29940975567778, 'addr_state_MA': 10.110592376611661, 'addr_state_AZ':
10.032653547892291, 'emp_type_Admin': 9.900472150565802, 'addr_state_WA':
9.604463414936513, 'addr_state_CO': 9.335793678183599, 'addr_state_MN':
```

8.09825045258343, 'addr_state_MO': 7.3337478164256416, 'addr_state_IN':
 7.156485816507435, 'addr_state_CT': 7.036377600696594, 'addr_state_TN':
 6.755168526348449, 'addr_state_NV': 6.5577667265598, 'emp_type_Manlab':
 6.5364938162900845, 'purpose_home_improvement': 6.521215768226548,
 'addr_state_WI': 6.175510070761902, 'addr_state_AL': 6.011851495010638,
 'emp_type_Clerk': 5.952461595533043, 'addr_state_OR': 5.87356020795903,
 'addr_state_SC': 5.763179271570182, 'addr_state_LA': 5.739137666645026,
 'purpose_other': 5.6022187340814265, 'addr_state_KY': 4.837516798054074,
 'addr_state_OK': 4.629869513098388, 'addr_state_KS': 4.55823347482431,
 'addr_state_AR': 3.9876202810025814, 'addr_state_UT': 3.8173769222531315,
 'emp_type_Designer': 3.29235512104278, 'addr_state_NM': 3.224544717816936,
 'addr_state_HI': 3.0593367943517693, 'addr_state_WV': 2.979129559660258,
 'addr_state_NH': 2.9350223110414437, 'purpose_major_purchase':
 2.8961488105820443, 'emp_type_Law': 2.8076997711159604, 'addr_state_RI':
 2.755806044102352, 'il_util': 2.7308005905563095, 'addr_state_MS':
 2.725067872453719, 'total_acc': 2.3412784187313673, 'emp_type_Clergy':
 2.29732842794935, 'inq_last_12m': 2.272432327564242, 'open_acc':
 2.236904562835468, 'purpose_small_business': 2.176687879808439, 'addr_state_MT':
 2.155466265716131, 'addr_state_DE': 2.134154484319103, 'addr_state_DC':
 2.1003725611804636, 'inq_fi': 2.027959693450723, 'open_acc_6m':
 1.9692331466328516, 'purpose_medical': 1.9542082569980301, 'addr_state_WY':
 1.916714690815398, 'grade': 1.8826323466447483, 'tot_cur_bal':
 1.8817770223027797, 'addr_state_SD': 1.8212586202969663, 'dti_joint':
 1.8138557594089295, 'addr_state_VT': 1.8123857317559084, 'loan_amnt':
 1.770450023225754, 'verification_status_Verified': 1.6876383832952806,
 'max_bal_bc': 1.6706128393200244, 'home_ownership_RENT': 1.6655753043101997,
 'purpose_moving': 1.6129660556591545, 'term': 1.6119942311081126,
 'mths_since_last_delinq': 1.6025300556111002,
 'verification_status_joint_Verified': 1.5951578280856389, 'purpose_vacation':
 1.5330136777572476, 'addr_state_NE': 1.5329177157986424,
 'verification_status_Source Verified': 1.5316244200874691,
 'mths_since_last_major_derog': 1.5014338644346426, 'revol_bal':
 1.4842587244820489, 'purpose_house': 1.4200368732846809, 'annual_inc':
 1.411788066763737, 'dti': 1.3580820245624448, 'revol_util': 1.3431387213056754,
 'total_cu_tl': 1.313277509616113, 'earliest_cr_line': 1.2774740226005623,
 'purpose_wedding': 1.2676206035373658, 'addr_state_ME': 1.2389826625507752,
 'emp_length': 1.23015208079886, 'addr_state_ND': 1.2176282318766618,
 'verification_status_joint_Source Verified': 1.21746162184719,
 'last_credit_pull_d': 1.2155761978077833, 'delinq_2yrs': 1.203056931316467,
 'inq_last_6mths': 1.1700827644410172, 'home_ownership_OWN': 1.1557264940598269,
 'last_pymnt_amnt': 1.1515113802060482, 'initial_list_status': 1.128468138127505,
 'pub_rec': 1.0729805349528614, 'purpose_renewable_energy': 1.0650906758641836,
 'purpose_educational': 1.0559243025772849, 'acc_now_delinq': 1.0219935281968442,
 'collections_12_mths_ex_med': 1.014027801963101, 'addr_state_IA':
 1.0075443124569217, 'addr_state_ID': 1.005791675253549, 'home_ownership_OTHER':
 1.0037427138737391, 'tot_coll_amt': 1.0008889405120973, 'pymnt_plan':
 1.0001515259612508}

```
[197]: corr_dict(X_scaled, corr = .5)
```

```
{'il_util': 4, 'inq_last_12m': 3, 'open_acc_6m': 2, 'inq-fi': 2,
'purpose_credit_card': 1, 'dti-joint': 1, 'mths_since_last_major_derog': 1,
'max_bal_bc': 1, 'total_acc': 1, 'verification_status_Source Verified': 1,
'verification_status_Verified': 1, 'verification_status_joint_Verified': 1,
'mths_since_last_delinq': 1, 'purpose_debt_consolidation': 1, 'open_acc': 1}
```

```
[198]: corr_list(corr= .6, min_corr = 1)
```

```
{'il_util': ['open_acc_6m', 'max_bal_bc'], 'inq_last_12m': ['inq-fi'],
'open_acc_6m': ['il_util'], 'inq-fi': ['inq_last_12m'], 'purpose_credit_card':
['purpose_debt_consolidation'], 'dti-joint': [], 'mths_since_last_major_derog':
[], 'max_bal_bc': ['il_util'], 'total_acc': ['open_acc'],
'verification_status_Source Verified': [], 'verification_status_Verified': [],
'verification_status_joint_Verified': [], 'mths_since_last_delinq': [],
'purpose_debt_consolidation': ['purpose_credit_card'], 'open_acc':
['total_acc']}
```

```
[195]: X_scaled.drop(['emp_type_other', 'addr_state_CA', ], axis = 1, inplace = True)
```

```
[196]: vif(X_scaled)
```

115 out of 121 features have a vif < 2.5

```
{'purpose_debt_consolidation': 24.96761119362853, 'purpose_credit_card':
18.93961785627865, 'purpose_home_improvement': 6.5212061819225475,
'purpose_other': 5.602217767774178, 'purpose_major_purchase': 2.896130530451369,
'il_util': 2.730798750721492, 'total_acc': 2.3412250411725206, 'inq_last_12m':
2.2724305415892028, 'open_acc': 2.2366887915277776, 'purpose_small_business':
2.1766772917631436, 'inq-fi': 2.027958161959543, 'open_acc_6m':
1.969223020722815, 'purpose_medical': 1.9541917795725594, 'grade':
1.882609031023419, 'tot_cur_bal': 1.8817109187706282, 'dti-joint':
1.8138544467538595, 'loan_amnt': 1.7700889155867419,
'verification_status_Verified': 1.6876008414952248, 'max_bal_bc':
1.67059263851143, 'home_ownership_RENT': 1.6649091693054299, 'purpose_moving':
1.612957560969563, 'term': 1.6119723224069213, 'mths_since_last_delinq':
1.6024830947155813, 'verification_status_joint_Verified': 1.5951568502705917,
'purpose_vacation': 1.5330036609346471, 'verification_status_Source Verified':
1.531580136517418, 'mths_since_last_major_derog': 1.501287811129522,
'revol_bal': 1.4842361387733702, 'addr_state_TX': 1.4710102981603153,
'addr_state_NY': 1.4451569706399565, 'purpose_house': 1.4200324341093116,
'annual_inc': 1.411719771951544, 'addr_state_FL': 1.3856905162976025,
'emp_type_Manager': 1.3846289802441962, 'emp_type_Unemployed':
1.375953539616218, 'dti': 1.3579245825891364, 'revol_util': 1.3429068081647426,
'total_cu_tl': 1.3132247928287366, 'emp_type_Technical': 1.291253306726926,
'earliest_cr_line': 1.2774538519434466, 'purpose_wedding': 1.2676155303887011,
'addr_state_IL': 1.2367942449247422, 'emp_length': 1.229640725680606,
'verification_status_joint_Source Verified': 1.2174567422552078,
'last_credit_pull_d': 1.2153732889718927, 'addr_state_PA': 1.2150862304066916,
```

```
'addr_state_OH': 1.2136184480397716, 'addr_state_NJ': 1.211756939901659,
'delinq_2yrs': 1.2030519679157148, 'addr_state_GA': 1.2027119102739965,
'addr_state_NC': 1.1744377442178562, 'addr_state_VA': 1.173599554052189,
'inq_last_6mths': 1.170005112914522, 'addr_state_MI': 1.167534979071377,
'emp_type_Healer': 1.1595810943404903, 'home_ownership_OWN': 1.1557259923910626,
'last_pymnt_amnt': 1.1515071804270436, 'emp_type_Executive': 1.1435806737986498,
'addr_state_MD': 1.1398098047832472, 'addr_state_AZ': 1.1366589905091007,
'addr_state_MA': 1.133407108180576, 'initial_list_status': 1.128389047723271,
'addr_state_WA': 1.127056734443586, 'addr_state_CO': 1.126097731275573,
'emp_type_Director': 1.1185216937061957, 'addr_state_MN': 1.111092139024878,
'addr_state_IN': 1.1080028479183648, 'addr_state_MO': 1.104552350953598,
'emp_type_Education': 1.0999518529880863, 'addr_state_TN': 1.0958630462702927,
'addr_state_CT': 1.0905781695634875, 'addr_state_AL': 1.08680622935313,
'emp_type_Vol': 1.0856810680418014, 'emp_type_Operator': 1.0853554300692798,
'addr_state_NV': 1.0841331667014709, 'addr_state_WI': 1.0838599850181452,
'addr_state_SC': 1.0790555747547514, 'addr_state_LA': 1.0789413239840797,
'emp_type_Analyst': 1.0748841656080348, 'pub_rec': 1.0729507946244885,
'addr_state_OR': 1.0727078127290053, 'emp_type_Service': 1.0708955288083002,
'emp_type_Assistant': 1.06592784508023, 'addr_state_KY': 1.0656549519342704,
'purpose_renewable_energy': 1.065090388165688, 'addr_state_OK':
1.0620347311872813, 'emp_type_Sales': 1.0590093693419453, 'addr_state_KS':
1.0584012271526348, 'purpose_educational': 1.0559240389158238, 'addr_state_AR':
1.0519412080590325, 'addr_state_UT': 1.0447937645886771, 'emp_type_Accountant':
1.0364811525268898, 'addr_state_NM': 1.0364600347939292, 'emp_type_Finance':
1.036374254004344, 'addr_state_WV': 1.035090185088152, 'emp_type_Admin':
1.0322020039686417, 'addr_state_MS': 1.0316747889622462, 'addr_state_HI':
1.0312558285234903, 'addr_state_NH': 1.0299181430529991, 'addr_state_RI':
1.0263658039473804, 'emp_type_Manlab': 1.0222894653635657, 'acc_now_delinq':
1.0219905270476908, 'emp_type_Clerk': 1.0206809614570438, 'addr_state_MT':
1.0188570977775266, 'addr_state_DE': 1.0178363363690994, 'addr_state_DC':
1.0169735661498065, 'addr_state_WY': 1.0158492688032048, 'addr_state_SD':
1.0146824568977704, 'collections_12_mths_ex_med': 1.014024859903138,
'addr_state_VT': 1.0135005534468404, 'addr_state_NE': 1.0093473303132385,
'emp_type_Designer': 1.0091182896478026, 'emp_type_Law': 1.007220539009543,
'emp_type_Clergy': 1.005854682864526, 'addr_state_ME': 1.0052699934125424,
'addr_state_ND': 1.0041975392687317, 'home_ownership_OTHER': 1.003742575348957,
'addr_state_IA': 1.0012945682069592, 'tot_coll_amt': 1.0008888908688367,
'addr_state_ID': 1.000435302095682, 'pymnt_plan': 1.0001514705171868}
```

```
[203]: X_scaled.drop(['purpose_debt_consolidation'], axis = 1, inplace = True)
```

```
[204]: vif(X_scaled)
```

```
119 out of 120 features have a vif < 2.5
{'il_util': 2.7307888242024836, 'total_acc': 2.34122363770942, 'inq_last_12m':
2.2724287817188364, 'open_acc': 2.2362154778845387, 'inq_fi': 2.027957789717889,
'open_acc_6m': 1.9692149833036792, 'grade': 1.8825266359843544, 'tot_cur_bal':
1.8816856234174637, 'dti_joint': 1.8138513300300452, 'loan_amnt':
```

1.76378376631423, 'verification_status_Verified': 1.6875362019651317,
'max_bal_bc': 1.6705891495577427, 'home_ownership_RENT': 1.6649046553273368,
'term': 1.6113500489848105, 'mths_since_last_delinq': 1.6024581712519677,
'verification_status_joint_Verified': 1.5951548391405679,
'verification_status_Source Verified': 1.5314865482589934,
'mths_since_last_major_derog': 1.5012617294420394, 'revol_bal':
1.4840021620392303, 'addr_state_TX': 1.471009052033995, 'addr_state_NY':
1.445139520674233, 'annual_inc': 1.4115003680493299, 'addr_state_FL':
1.3856881112765589, 'emp_type_Manager': 1.3845719470330107,
'emp_type_Unemployed': 1.3758843463096149, 'dti': 1.3568266538343439,
'revol_util': 1.3404689847207572, 'total_cu_tl': 1.313216696550245,
'emp_type_Technical': 1.291252864896083, 'earliest_cr_line': 1.2774538426875484,
'addr_state_IL': 1.236793043311825, 'emp_length': 1.2295347307952664,
'verification_status_joint_Source Verified': 1.2174567012313438,
'addr_state_PA': 1.2150812972677942, 'last_credit_pull_d': 1.2140420349666088,
'addr_state_OH': 1.2136099740784634, 'addr_state_NJ': 1.2117561619082107,
'delinq_2yrs': 1.2030420768765697, 'addr_state_GA': 1.2027109402714755,
'addr_state_NC': 1.1744360565205412, 'addr_state_VA': 1.1735991837449828,
'inq_last_6mths': 1.1699938360933184, 'addr_state_MI': 1.1675327406173073,
'emp_type_Healer': 1.1595726960576076, 'home_ownership_OWN': 1.1556346606602725,
'last_pymnt_amnt': 1.1513341738552931, 'emp_type_Executive': 1.1435483886317626,
'addr_state_MD': 1.139809707927233, 'addr_state_AZ': 1.1366565574300957,
'addr_state_MA': 1.1333957445531833, 'initial_list_status': 1.1282615470377386,
'addr_state_WA': 1.1270557903662477, 'addr_state_CO': 1.1260946762503712,
'purpose_credit_card': 1.1221073903675312, 'emp_type_Director':
1.1184903635699266, 'addr_state_MN': 1.111089812879824, 'addr_state_IN':
1.1079998719019568, 'addr_state_MO': 1.1045512858739144, 'emp_type_Education':
1.099945224104737, 'addr_state_TN': 1.0958594251080822, 'addr_state_CT':
1.0905774354390236, 'addr_state_AL': 1.0868059692953485,
'purpose_home_improvement': 1.0859889095512643, 'emp_type_Vol':
1.0856756631964963, 'emp_type_Operator': 1.0853548688948504, 'addr_state_NV':
1.0841331538525212, 'addr_state_WI': 1.0838585857943173, 'purpose_other':
1.0792270558935455, 'addr_state_SC': 1.0790555482931103, 'addr_state_LA':
1.0789353312786993, 'emp_type_Analyst': 1.0748773175088338, 'pub_rec':
1.072882602580039, 'addr_state_OR': 1.072704089026199, 'emp_type_Service':
1.0708780679243441, 'emp_type_Assistant': 1.065852142672351, 'addr_state_KY':
1.0656546322556582, 'addr_state_OK': 1.062033248172776, 'emp_type_Sales':
1.0590011240943042, 'addr_state_KS': 1.058401215690259, 'addr_state_AR':
1.051938265638769, 'addr_state_UT': 1.0447901011894005,
'purpose_small_business': 1.0400080197076909, 'emp_type_Accountant':
1.0364722234800476, 'addr_state_NM': 1.0364568579215143, 'emp_type_Finance':
1.0363311115149034, 'addr_state_WV': 1.0350897034065472, 'emp_type_Admin':
1.0321939421946484, 'addr_state_MS': 1.0316741156182487, 'addr_state_HI':
1.0312541185818314, 'addr_state_NH': 1.0299139919431572,
'purpose_major_purchase': 1.029421341791009, 'addr_state_RI': 1.02636078843678,
'emp_type_Manlab': 1.0222840437021758, 'acc_now_delinq': 1.0219903082614925,
'emp_type_Clerk': 1.0206739950427475, 'purpose_moving': 1.019002022618076,
'addr_state_MT': 1.01885669067497, 'purpose_medical': 1.0178710196559664,

```
'addr_state_DE': 1.0178345137750904, 'addr_state_DC': 1.0169695912928167,
'addr_state_WY': 1.0158399393098407, 'addr_state_SD': 1.0146792946358334,
'collections_12_mths_ex_med': 1.0140236676392183, 'purpose_vacation':
1.0138077193121813, 'addr_state_VT': 1.0135003335191375, 'purpose_wedding':
1.012731673605749, 'purpose_educational': 1.0115037506358027, 'purpose_house':
1.010585488111, 'addr_state_NE': 1.0093460943965555, 'emp_type_Designer':
1.0091124645236929, 'emp_type_Law': 1.0072052221037187, 'emp_type_Clergy':
1.005846092220493, 'addr_state_ME': 1.00526910748636, 'addr_state_ND':
1.0041965609613241, 'home_ownership_OTHER': 1.003740834803285,
'purpose_renewable_energy': 1.0016474231528574, 'addr_state_IA':
1.0012934740261756, 'tot_coll_amt': 1.0008888906939468, 'addr_state_ID':
1.0004330794741287, 'pymnt_plan': 1.000151470516797}
```

While I wanted all of the features to be under 2.5, The dataset is still large and running the vif function takes over 10 minutes to run. I will first perform feature selection and then I will return to vif.

```
[206]: y_binary = np.where(loan['loan_status'].isin(['Fully Paid', 'Current', 'Issued']), 1, 0)
```

```
##
```

Lasso Regularization

Lasso Regularization is a technique for variable selection that uses regression to evaluate the effect that features have on a target variable. The idea is to add a penalty term that contains the coefficient or slope of the variables with respect to the target variable multiplied by parameter.

$$Loss(\beta_1, \dots, \beta_n) = SSD + \alpha |\beta|$$

Where β is the coefficient of the feature, SSD is the sum of squared distances of the point to the regression line and α is the penalty's parameter. If there is more than a single feature, then we can sum the coefficients:

$$Loss(\beta_1, \dots, \beta_n) = SSD + \alpha * \sum_{i=1}^n |\beta_i|$$

Lasso aims to reduce the loss in equation (3). If a feature is not important, then changing its slope will not move the regression line close enough to the data points to decrease the loss function with a non-zero value. If a feature is important, the regression line will move towards the data and minimize the SSD faster than it increases the penalty term. Since only features with non-zero coefficients are meaningful, I can discard all the features with a coefficient of zero.

First we will run a gridsearch to find the best value of α , and then we will use the coefficients from that value of α to determine which features stay in the model.

```
[207]: # Create Lasso model
lasso = Lasso(max_iter = 50000)

# Define hyperparameter grid with a value less than 00.5 since that was the best
params = {'alpha': np.linspace(.0001, .01, 20)}

# Perform grid search
```

```

grid_search = GridSearchCV(estimator=lasso, param_grid=params, cv=8)

# fit the gridsearch of parameters to the data
grid_search.fit(X_scaled, y_binary)

# Print best hyperparameters
print("Best hyperparameters: ", grid_search.best_params_)

```

Best hyperparameters: {'alpha': 0.01}

```

[208]: # Find the coefficients of lasso regularization
lasso = Lasso(alpha = .001)
# Fit the lasso regularization to the data
lasso.fit(X_scaled, y_binary)
# Create a dictionary of all the features and their corresponding lasso
↪ coefficients
lasso_dict = {X_scaled.columns[i]:lasso.coef_[i] for i in range(len(X_scaled.
↪ columns)) if list(lasso.coef_)[i] != 0}
# Create a list of all the features that don't have a lasso coefficient of zero
lasso_features = [X_scaled.columns[i] for i in range(len(X_scaled.columns)) if
↪ list(lasso.coef_)[i] != 0]
# Print the features and coefficients
print(lasso_dict)
# Print the number of features are left from lasso
print(len(lasso_dict))

```

```

{'loan_amnt': -0.010391316872062236, 'term': 0.002948502470419479, 'grade':
0.03788254325234942, 'annual_inc': 0.0022825117307152722, 'pymnt_plan':
-0.00023065455231244378, 'inq_last_6mths': -0.016965471271241, 'pub_rec':
0.00328094929656237, 'revol_bal': 0.00017543217177538116, 'revol_util':
-0.005030065433174285, 'total_acc': -0.001725540484437378,
'initial_list_status': 0.010934774737318688, 'last_pymnt_amnt':
0.04433076049762753, 'last_credit_pull_d': 0.051610620607713674,
'collections_12_mths_ex_med': 0.0004269000786191071,
'mths_since_last_major_derog': -0.003163609410562672, 'dti_joint':
0.0007302915041974671, 'tot_cur_bal': 0.0021794962954110103, 'open_acc_6m':
0.0023301194338613833, 'il_util': 0.004372566069389636, 'max_bal_bc':
0.0015211874387965623, 'inq_last_12m': 0.001700906513944068,
'emp_type_Accountant': 0.0014560825330323078, 'emp_type_Admin':
0.0011180935599351067, 'emp_type_Analyst': 0.0020839310189119495,
'emp_type_Assistant': 0.001716231097294549, 'emp_type_Clergy':
0.0004086759716117618, 'emp_type_Clerk': 0.0003684051643701604,
'emp_type_Designer': 8.386470181023748e-05, 'emp_type_Director':
0.003250926563873108, 'emp_type_Education': 0.001905001710557324,
'emp_type_Executive': 0.0028161364298661838, 'emp_type_Healer':
0.001802276485713458, 'emp_type_Manager': 0.00534780536249633,
'emp_type_Operator': 0.00017582019867604627, 'emp_type_Technical':
0.0017201398884130043, 'emp_type_Vol': 0.0013373643692269049,

```



```
'home_ownership_OTHER': -9.787280568326191e-05, 'home_ownership_RENT':
-0.0034307500430190156, 'verification_status_Source Verified':
0.0006960764915860841, 'verification_status_Verified': -0.002840233859866113,
'purpose_credit_card': 0.0010696145226728286, 'purpose_major_purchase':
0.0006338286983715384, 'purpose_other': 0.00030442720791727033,
'purpose_small_business': -0.0025923711676408053, 'purpose_wedding':
0.0008600910780396444, 'addr_state_AL': -2.717313790854032e-05, 'addr_state_CO':
0.00023737701126124214, 'addr_state_DC': 0.0003228253473749971, 'addr_state_FL':
-0.0008238996713227565, 'addr_state_IL': 0.0018622525159464468, 'addr_state_KS':
0.0003714578370217212, 'addr_state_ME': 0.0005063581439605373, 'addr_state_MS':
0.0004445402559083527, 'addr_state_ND': 0.00031819850705387936, 'addr_state_NE':
0.0007900323376744395, 'addr_state_NH': 1.8891369076422682e-05, 'addr_state_NV':
-0.0006201602888351354, 'addr_state_NY': -0.000804910934598037, 'addr_state_SC':
0.0004239016727849242, 'addr_state_TX': 0.0011805290982262464, 'addr_state_VA':
-0.00022906652147532304}
61
```

```
[211]: df = X_scaled[lasso_dict.keys()]
```

```
[212]: vif(df)
```

```
60 out of 61 features have a vif < 2.5
{'il_util': 2.503648536830864, 'open_acc_6m': 1.9535214353509527, 'tot_cur_bal':
1.770654329297814, 'grade': 1.7439133806542608, 'loan_amnt': 1.7138957829652597,
'max_bal_bc': 1.6615584854813974, 'inq_last_12m': 1.6414880361328827,
'verification_status_Verified': 1.625458267723455, 'term': 1.5872880178470525,
'verification_status_Source Verified': 1.523078231850744, 'revol_bal':
1.4016325218077512, 'annual_inc': 1.330297515878795, 'home_ownership_RENT':
1.3073547671221735, 'total_acc': 1.2448280761753376, 'revol_util':
1.2440126800569045, 'emp_type_Manager': 1.2293581030145195,
'last_credit_pull_d': 1.1823828739055637, 'emp_type_Technical':
1.1632605142157963, 'inq_last_6mths': 1.1554307323054438, 'last_pymnt_amnt':
1.140772579699193, 'initial_list_status': 1.1210513382502567,
'emp_type_Executive': 1.1074575667950384, 'emp_type_Healer': 1.0974140074627614,
'purpose_credit_card': 1.090279220777798, 'mths_since_last_major_derog':
1.0872113114911621, 'emp_type_Director': 1.0761671055142379,
'emp_type_Education': 1.0593520442957665, 'purpose_other': 1.0584937234060041,
'addr_state_NY': 1.0554819306031729, 'addr_state_TX': 1.0515488522453607,
'emp_type_Vol': 1.0511714498170022, 'emp_type_Operator': 1.0502161947572395,
'pub_rec': 1.0497176483600994, 'emp_type_Analyst': 1.0460144044523427,
'emp_type_Assistant': 1.0392069571630331, 'addr_state_FL': 1.0390726711554674,
'purpose_small_business': 1.0315205970379795, 'addr_state_IL':
1.0262364921930622, 'emp_type_Accountant': 1.0217594682580176,
'purpose_major_purchase': 1.0206314648766086, 'addr_state_VA':
1.020258725151585, 'emp_type_Admin': 1.0192132595134875, 'addr_state_CO':
1.0133330322712164, 'emp_type_Clerk': 1.0117102129881905, 'addr_state_AL':
1.0112905135137134, 'collections_12_mths_ex_med': 1.010775638772582,
'purpose_wedding': 1.0107579856757944, 'addr_state_NV': 1.0095052605905286,
```



```
'addr_state_SC': 1.0094403966131993, 'addr_state_KS': 1.0067548263507855,  
'emp_type_Designer': 1.0054892655608172, 'addr_state_MS': 1.005021981570084,  
'emp_type_Clergy': 1.0034703330142034, 'addr_state_DC': 1.0034678554709804,  
'addr_state_NH': 1.0032999442711354, 'dti_joint': 1.0031683280406816,  
'home_ownership_OTHER': 1.0030594050709054, 'addr_state_NE': 1.0015282820188407,  
'addr_state_ME': 1.0014568397523846, 'addr_state_ND': 1.0009243078467163,  
'pymnt_plan': 1.000111548764211}
```

```
[213]: df.to_csv('./data/interim/wrangled', index = False)
```